

Spotify_EDA_v4 (1) (1)

April 11, 2023

1 Spotify Machine Learning

In this PIC 16A Final project, our group decided to use the Spotify Dataset sourced from [Kaggle](#) in order to see if we can accurately predict the characteristics of a song that would make it to the top of the Spotify Top 200 Charts.

2 Part 1. Group Contributions Statement

All three of us wrote the data acquisition and preparation. Liyuan wrote the functions used to clean and split the data and Sidney also contributed to the functions and manual recoding. Jia Shing led Figure 1 and the K-Neighbor model and the Random Forest model. LiYuan led Figure 3 and 4 and the SVM model in addition to working on Random Forest. Jia Shing also wrote the functions used to ultimately graph the models. Sidney led Figure 5, Table 1, and the Logistic Regression model. Each person wrote the explanations of their own figures and models. Sidney wrote the conclusion.

We all checked each other's work and made revisions to code and writing

3 Part 2: Data Import and Cleaning

3.0.1 Part 2.1: Clean Datapoints for EDA

In this section of the notebook, we wish to achieve the following: - Explore all the variables available in the dataset - Clean the dataset - Check for NAs/NULL values (Remove columns with empty data) - Remove duplicates - Check DType and ensure they are correct

We decided to clean our data set **before** splitting into test and train because we are working with a new dataset and wanted to make sure that we could thoroughly explore all of the data well enough to know how to remove outliers. This is because simply splitting the dataset into training and testing sets may result in one or more classes being severely underrepresented in one or both sets, which could negatively impact the model's accuracy. Additionally, due to some of the aspects of our model, we needed to conduct manual inspection/use human judgment, which requires a thorough understanding of all the context of the data.

Our "cleaning" was strictly limited to mapping or dropping non compatible data types, we only did labelEncoding/scaling after splitage, not in this step. Although this step may have sacrificed some accuracy, we believe that it was ultimately more important to have a better visualization of our dataset.

```
[1]: # Importing some basic libraries
import pandas as pd
import numpy as np
```

```
[2]: # Importing the dataset
url = "https://raw.githubusercontent.com/jia-shing/pic16a-spotify/main/
      ↳spotify_dataset.csv"
spotify= pd.read_csv(url)
spotify
```

```
[2]:
```

	Index	Highest Charting Position	Number of Times Charted \
0	1	1	8
1	2	2	3
2	3	1	11
3	4	3	5
4	5	5	1
...
1551	1552	195	1
1552	1553	196	1
1553	1554	197	1
1554	1555	198	1
1555	1556	199	1

	Week of Highest Charting	Song Name	Streams \
0	2021-07-23--2021-07-30	Beggin'	48,633,449
1	2021-07-23--2021-07-30	STAY (with Justin Bieber)	47,248,719
2	2021-06-25--2021-07-02	good 4 u	40,162,559
3	2021-07-02--2021-07-09	Bad Habits	37,799,456
4	2021-07-23--2021-07-30	INDUSTRY BABY (feat. Jack Harlow)	33,948,454
...
1551	2019-12-27--2020-01-03	New Rules	4,630,675
1552	2019-12-27--2020-01-03	Cheirosa - Ao Vivo	4,623,030
1553	2019-12-27--2020-01-03	Havana (feat. Young Thug)	4,620,876
1554	2019-12-27--2020-01-03	Surtada - Remix Brega Funk	4,607,385
1555	2019-12-27--2020-01-03	Lover (Remix) [feat. Shawn Mendes]	4,595,450

	Artist	Artist Followers	Song ID \
0	Måneskin	3377762.0	3Wrjm47oTz2sjIgck1115e
1	The Kid LAROI	2230022.0	5HCyWlXZPP0y6Gqq8TgA20
2	Olivia Rodrigo	6266514.0	4ZtFanR9U6ndgddUvNcjcG
3	Ed Sheeran	83293380.0	6PQ88X9TkUIAUIZJHW2upE
4	Lil Nas X	5473565.0	27NovPIUIRrOZoCHxABJwK
...
1551	Dua Lipa	27167675.0	2ekn2ttSfGqwhhate0LSR0
1552	Jorge & Mateus	15019109.0	2PWjKmJyTZEDpmOUa3a5da
1553	Camila Cabello	22698747.0	1rfofaqEpACxVEHIZBJe6W
1554	Dadá Boladão, Tati Zaqui, OIK	208630.0	5F8ffc8KWKNAwllr5WsW0r

1555 Taylor Swift 42227614.0 3i9UVldZ0E0aD0JnyfAZZ0

	Genre	Danceability
0	['indie rock italiano', 'italian pop']	0.714
1	['australian hip hop']	0.591
2	['pop']	0.563
3	['pop', 'uk pop']	0.808
4	['lgbtq+ hip hop', 'pop rap']	0.736
...
1551	['dance pop', 'pop', 'uk pop']	0.762
1552	['sertanejo', 'sertanejo universitario']	0.528
1553	['dance pop', 'electropop', 'pop', 'post-teen ...	0.765
1554	['brega funk', 'funk carioca']	0.832
1555	['pop', 'post-teen pop']	0.448

	Energy	Loudness	Speechiness	Acousticness	Liveness	Tempo
0	0.800	-4.808	0.0504	0.12700	0.3590	134.002
1	0.764	-5.484	0.0483	0.03830	0.1030	169.928
2	0.664	-5.044	0.1540	0.33500	0.0849	166.928
3	0.897	-3.712	0.0348	0.04690	0.3640	126.026
4	0.704	-7.409	0.0615	0.02030	0.0501	149.995
...
1551	0.700	-6.021	0.0694	0.00261	0.1530	116.073
1552	0.870	-3.123	0.0851	0.24000	0.3330	152.370
1553	0.523	-4.333	0.0300	0.18400	0.1320	104.988
1554	0.550	-7.026	0.0587	0.24900	0.1820	154.064
1555	0.603	-7.176	0.0640	0.43300	0.0862	205.272

	Duration (ms)	Valence	Chord
0	211560.0	0.589	B
1	141806.0	0.478	C#/Db
2	178147.0	0.688	A
3	231041.0	0.591	B
4	212000.0	0.894	D#/Eb
...
1551	209320.0	0.608	A
1552	181930.0	0.714	B
1553	217307.0	0.394	D
1554	152784.0	0.881	F
1555	221307.0	0.422	G

[1556 rows x 23 columns]

```
[3]: # Exploring some basic metadata about the dataset
spotify.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 1556 entries, 0 to 1555

Data columns (total 23 columns):

#	Column	Non-Null Count	Dtype
0	Index	1556 non-null	int64
1	Highest Charting Position	1556 non-null	int64
2	Number of Times Charted	1556 non-null	int64
3	Week of Highest Charting	1556 non-null	object
4	Song Name	1556 non-null	object
5	Streams	1556 non-null	object
6	Artist	1556 non-null	object
7	Artist Followers	1545 non-null	float64
8	Song ID	1545 non-null	object
9	Genre	1545 non-null	object
10	Release Date	1545 non-null	object
11	Weeks Charted	1556 non-null	object
12	Popularity	1545 non-null	float64
13	Danceability	1545 non-null	float64
14	Energy	1545 non-null	float64
15	Loudness	1545 non-null	float64
16	Speechiness	1545 non-null	float64
17	Acousticness	1545 non-null	float64
18	Liveness	1545 non-null	float64
19	Tempo	1545 non-null	float64
20	Duration (ms)	1545 non-null	float64
21	Valence	1545 non-null	float64
22	Chord	1545 non-null	object

dtypes: float64(11), int64(3), object(9)

memory usage: 279.7+ KB

From here, we can see that we have 1556 non-null values in certain fields and 1545 in others. However, we note that the *DType* of certain fields such as “Week of Highest Charting”, “Streams”, “Song ID” etc are listed as *object* instead of *float* or *int*, which is what we expect from observing the dataset.

Upon closer inspection of the dataset, we see that instead of a NULL value, there were some cells that have a string with just a space. Therefore, all our numeric values are coerced into strings, which makes our *DType* an *object*.

To rectify this, we will use *pandas* to replace these blanks with *NaN* so that we can drop them.

```
[4]: # Replacing the empty values with NaN
spotify = spotify.replace(to_replace=" ", value=np.nan)
spotify.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 1556 entries, 0 to 1555

Data columns (total 23 columns):

#	Column	Non-Null Count	Dtype
---	--------	----------------	-------

0	Index	1556 non-null	int64
1	Highest Charting Position	1556 non-null	int64
2	Number of Times Charted	1556 non-null	int64
3	Week of Highest Charting	1556 non-null	object
4	Song Name	1556 non-null	object
5	Streams	1556 non-null	object
6	Artist	1556 non-null	object
7	Artist Followers	1545 non-null	float64
8	Song ID	1545 non-null	object
9	Genre	1545 non-null	object
10	Release Date	1545 non-null	object
11	Weeks Charted	1556 non-null	object
12	Popularity	1545 non-null	float64
13	Danceability	1545 non-null	float64
14	Energy	1545 non-null	float64
15	Loudness	1545 non-null	float64
16	Speechiness	1545 non-null	float64
17	Acousticness	1545 non-null	float64
18	Liveness	1545 non-null	float64
19	Tempo	1545 non-null	float64
20	Duration (ms)	1545 non-null	float64
21	Valence	1545 non-null	float64
22	Chord	1545 non-null	object

dtypes: float64(11), int64(3), object(9)

memory usage: 279.7+ KB

```
[5]: # Dropping the NaN values
spotify.dropna(inplace=True)

# resetting the index
spotify.reset_index(drop=True, inplace=True)

spotify.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1545 entries, 0 to 1544
Data columns (total 23 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Index                                1545 non-null   int64
1   Highest Charting Position            1545 non-null   int64
2   Number of Times Charted              1545 non-null   int64
3   Week of Highest Charting             1545 non-null   object
4   Song Name                           1545 non-null   object
5   Streams                             1545 non-null   object
6   Artist                              1545 non-null   object
7   Artist Followers                     1545 non-null   float64
```

```

8   Song ID          1545 non-null  object
9   Genre            1545 non-null  object
10  Release Date     1545 non-null  object
11  Weeks Charted    1545 non-null  object
12  Popularity       1545 non-null  float64
13  Danceability     1545 non-null  float64
14  Energy           1545 non-null  float64
15  Loudness         1545 non-null  float64
16  Speechiness      1545 non-null  float64
17  Acousticness     1545 non-null  float64
18  Liveness         1545 non-null  float64
19  Tempo            1545 non-null  float64
20  Duration (ms)    1545 non-null  float64
21  Valence          1545 non-null  float64
22  Chord            1545 non-null  object
dtypes: float64(11), int64(3), object(9)
memory usage: 277.7+ KB

```

We also noticed that for `streams`, the numbers in the dataframe were represented with commas in it i.e '7,234,437' format and we cannot use it for any numeric operation or plotting. So, the commas were removed and then it's data type was changed.

```

[6]: # Modifying and correcting the data type of the Streams field
spotify['Streams'] = spotify['Streams'].str.replace(',', '').apply(pd.
    ↳to_numeric)

spotify.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1545 entries, 0 to 1544
Data columns (total 23 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Index               1545 non-null  int64
1   Highest Charting Position 1545 non-null  int64
2   Number of Times Charted  1545 non-null  int64
3   Week of Highest Charting  1545 non-null  object
4   Song Name           1545 non-null  object
5   Streams             1545 non-null  int64
6   Artist              1545 non-null  object
7   Artist Followers     1545 non-null  float64
8   Song ID             1545 non-null  object
9   Genre               1545 non-null  object
10  Release Date        1545 non-null  object
11  Weeks Charted       1545 non-null  object
12  Popularity          1545 non-null  float64
13  Danceability        1545 non-null  float64
14  Energy              1545 non-null  float64

```

```

15 Loudness                1545 non-null    float64
16 Speechiness             1545 non-null    float64
17 Acousticness            1545 non-null    float64
18 Liveness                1545 non-null    float64
19 Tempo                   1545 non-null    float64
20 Duration (ms)          1545 non-null    float64
21 Valence                 1545 non-null    float64
22 Chord                   1545 non-null    object
dtypes: float64(11), int64(4), object(8)
memory usage: 277.7+ KB

```

Lastly, we would like to split up the `Week of Highest Charting` into a `datetime64[ns]` DType so that we can accurately calculate the number of days that a particular song was at its highest charting position.

Further, we would like to split this up into 2 columns, which is its start and end. This allows us to have the option of running some additional analysis and creating some interesting plots to see if a song's popularity is affected by the week of the year.

Side note: This is motivated by our observation of a particular genre called **adult standards**, which upon closer inspection of the dataset, was essentially Christmas music.

```

[7]: # Extracting the start and end week of the highest charting
spotify[['Week of Highest Charting_Start', 'Week of Highest Charting_End']] =
    ↳spotify['Week of Highest Charting'].str.split('--', expand=True)

# Changing column types to datetime
spotify['Week of Highest Charting_Start'] = pd.to_datetime(spotify['Week of
    ↳Highest Charting_Start'], yearfirst=True)
spotify['Week of Highest Charting_End'] = pd.to_datetime(spotify['Week of
    ↳Highest Charting_End'], yearfirst=True)
spotify.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1545 entries, 0 to 1544
Data columns (total 25 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Index                                1545 non-null   int64
1   Highest Charting Position            1545 non-null   int64
2   Number of Times Charted              1545 non-null   int64
3   Week of Highest Charting             1545 non-null   object
4   Song Name                           1545 non-null   object
5   Streams                             1545 non-null   int64
6   Artist                              1545 non-null   object
7   Artist Followers                     1545 non-null   float64
8   Song ID                             1545 non-null   object
9   Genre                               1545 non-null   object
10  Release Date                        1545 non-null   object

```

```

11 Weeks Charted          1545 non-null    object
12 Popularity             1545 non-null    float64
13 Danceability           1545 non-null    float64
14 Energy                 1545 non-null    float64
15 Loudness               1545 non-null    float64
16 Speechiness            1545 non-null    float64
17 Acousticness           1545 non-null    float64
18 Liveness               1545 non-null    float64
19 Tempo                  1545 non-null    float64
20 Duration (ms)         1545 non-null    float64
21 Valence                1545 non-null    float64
22 Chord                  1545 non-null    object
23 Week of Highest Charting_Start 1545 non-null    datetime64[ns]
24 Week of Highest Charting_End   1545 non-null    datetime64[ns]
dtypes: datetime64[ns](2), float64(11), int64(4), object(8)
memory usage: 301.9+ KB

```

3.0.2 Part 2.2: Remapping the Genre Column

From an initial observation, we can see that under the “Genre” field, we have **strings** that resembles a Python list, which contains the various genres of the particular song. From the output below, we also observe that the dataset seem to be extremely specific about the song genre. For instance, in `Genre[1551]`, we have `['dance pop', 'pop', 'uk pop']`.

For the purposes of our project, this level of specificity is unnecessary and would be difficult to use such specific genres to make more general analyses and conclusions about the performance of a song. We decided to narrow the genres down to 4 categories: rap, pop, r&b, and rock. Our reasoning behind this is that those are 4 relatively distinct genres of music that songs could be more easily classified into one box as.

We believe that this generalization will not adversely affect our performance, and would greatly simplify our EDA process; despite more specific subgenres a song may have, it is most likely that it could be associated with one of the following genres as an input to predict its trends

```
[8]: spotify['Genre']
```

```

[8]: 0          ['indie rock italiano', 'italian pop']
     1          ['australian hip hop']
     2          ['pop']
     3          ['pop', 'uk pop']
     4          ['lgbtq+ hip hop', 'pop rap']
     ...
1540          ['dance pop', 'pop', 'uk pop']
1541          ['sertanejo', 'sertanejo universitario']
1542          ['dance pop', 'electropop', 'pop', 'post-teen ...
1543          ['brega funk', 'funk carioca']
1544          ['pop', 'post-teen pop']
Name: Genre, Length: 1545, dtype: object

```


From the given `Genre` Series, we would like to split each string values into a list and combine them into a larger list. We would then use the `set` function to determine the total number of unique genres we have.

```
[9]: def delister(data):  
    """Takes in a single-column dataframe and split the string values in the_  
    ↪column into a list  
    Args:  
        data: Panda Series with values of type `string`  
  
    Returns:  
        data: Panda Series with values of type `list`  
    """  
  
    # Splitting the string values at single quotes "'"  
    data = data.apply(lambda x: x.split("'"))  
  
    # Appending the element of the list to a new list if the length of the_  
    ↪element is greater than 2 (i.e. removing the ",", "[" and "]" from the_  
    ↪list)  
    data = data.apply(lambda x: list(set([y for y in x if len(y) > 2])))  
  
    # Returning the series  
    return data  
  
delisted = delister(spotify['Genre'])  
delisted
```

```
[9]: 0          [italian pop, indie rock italiano]  
     1          [australian hip hop]  
     2          [pop]  
     3          [pop, uk pop]  
     4          [lgbtq+ hip hop, pop rap]  
     ...  
    1540          [pop, uk pop, dance pop]  
    1541          [sertanejo, sertanejo universitario]  
    1542          [pop, electropop, dance pop, post-teen pop]  
    1543          [funk carioca, brega funk]  
    1544          [pop, post-teen pop]  
    Name: Genre, Length: 1545, dtype: object
```

```
[10]: # A list comprehension to combine all the lists into a single list  
       combined = [i for x in delisted for i in x]  
  
       # Removing the duplicates from the list using the set function  
       unique_genres = list(set(combined))  
       print(unique_genres)
```

```
print("\n")

print(f"Number of unique types of genres: {len(unique_genres)}")
```

```
['brostep', 'german techno', 'pop rock', 'oakland hip hop', 'british soul',
'aussietronica', 'southern soul', 'viral pop', 'reggaeton colombiano', 'classic
rock', 'hard rock', 'indie pop', 'trap latino', 'bedroom pop', 'soul',
'swing', 'grime', 'german trap', 'eau claire indie', 'trap argentino', 'trap
boricua', 'torch song', 'dmv rap', 'italian pop rock', 'rap francais', 'uk
alternative hip hop', 'alternative r&b', 'mexican hip hop', 'pop nacional',
'francoton', 'rap metal', 'rhode island rap', 'etherpop', 'french hip hop',
'deep house', 'pacific islands pop', 'uk funky', 'electronic trap', 'jazz pop',
'motown', 'ohio hip hop', 'show tunes', 'funk ostentacao', 'irish singer-
songwriter', 'old school rap francais', 'lgbtq+ hip hop', 'deep german hip hop',
'hip pop', 'scandipop', 'afroswing', 'ranchera', 'rap', 'canadian trap', 'german
alternative rap', 'hardcore hip hop', 'lounge', 'mellow gold', 'nu metal',
'contemporary country', 'uk alternative pop', 'melodic rap', 'puerto rican pop',
'punk', 'slap house', 'london rap', 'pagode baiano', 'frauenrap', 'norteno',
'neo mellow', 'metalcore', 'regional mexican', 'rap belge', 'indie rockism',
'modern indie pop', 'chicago rap', 'new wave pop', 'pop argentino', 'soft rock',
'icelandic pop', 'art rock', 'german hip hop', 'progressive electro house',
'cumbia pop', 'colombian pop', 'southern hip hop', 'dream smp', 'latin hip hop',
'trance', 'post-grunge', 'alternative metal', 'sad rap', 'tekk', 'cubaton',
'jazz funk', 'swedish pop', 'k-pop', 'vancouver indie', 'modern rock', 'electro
latino', 'musical advocacy', 'moombahton', 'rockabilly', 'britpop', 'west coast
rap', 'florida rap', 'italian adult pop', 'reggaeton flow', 'nyc pop',
'heartland rock', 'basshall', 'forro', 'comic', 'a cappella', 'country rock',
'easy listening', 'alt z', 'rock-and-roll', 'hamburg hip hop', 'piano rock',
'permanent wave', 'conscious hip hop', 'german dance', 'k-pop girl group', 'rap
dominicano', 'melanesian pop', 'canadian contemporary r&b', 'afrofuturism',
'latin', 'uk dance', 'bubblegrunge', 'pop soul', 'brazilian hip hop', 'norwegian
pop', 'soundtrack', 'dance pop', 'celtic', 'canadian pop', 'turkish trap', 'trap
brasileiro', 'talent show', 'indie pop rap', 'deep euro house', 'miami hip hop',
'chicago soul', 'pop urbaine', 'neo soul', 'classic soul', 'country road',
'irish folk', 'art pop', 'post-teen pop', 'german underground rap', 'toronto
rap', 'jawaiian', 'nouvelle chanson francaise', 'europop', 'sunlensk tonlist',
'dominican pop', 'memphis hip hop', 'acoustic pop', 'glam rock', 'american folk
revival', 'australian rock', 'canadian hip hop', 'reggaeton', 'dembow', 'rap
conciencia', 'png pop', 'viral rap', 'nuevo regional mexicano', 'folk-pop', 'dfw
rap', 'classic uk pop', 'electropop', 'hip hop', 'social media pop', 'tennessee
hip hop', 'dreamo', 'new orleans rap', 'sertanejo', 'synthpop', 'indie pop',
'german pop', 'atl hip hop', 'minnesota hip hop', 'philly rap', 'kentucky hip
hop', 'yacht rock', 'funk 150 bpm', 'rap latina', 'dutch edm', 'venezuelan hip
hop', 'urban contemporary', 'italian indie pop', 'trancecore', 'houston rap',
'australian pop', 'italian hip hop', 'champeta', 'brega funk', 'funk pop',
'k-rap', 'detroit hip hop', 'rave funk', 'perreo', 'trap triste', 'modern
```

alternative rock', 'big room', 'east coast hip hop', 'adult standards', 'pop house', 'pop edm', 'indietronica', 'new french touch', 'pop reggaeton', 'chill r&b', 'uk pop', 'nz pop', 'country', 'north carolina hip hop', 'garage rock', 'sertanejo universitario', 'drill', 'italian pop', 'corrido', 'indie rock italiano', 'dutch pop', 'melodic metalcore', 'edm', 'funk paulista', 'funk carioca', 'pittsburgh rap', 'uk metalcore', 'pop r&b', 'cali rap', 'indonesian pop', 'deep underground hip hop', 'mariachi', 'indie cafe pop', 'shiver pop', 'r&b en espanol', 'quiet storm', 'south african house', 'german cloud rap', 'trap', 'boston hip hop', 'panamanian pop', 'queens hip hop', 'atl trap', 'hollywood', 'vapor trap', 'indie surf', 'progressive house', 'albanian hip hop', 'latin viral pop', 'alternative pop rock', 'beatlesque', 'canadian latin', 'r&b', 'plugg', 'indie r&b', 'surf punk', 'folktronica', 'vegas indie', 'meme rap', 'uk hip hop', 'underground hip hop', 'sudanese pop', 'pop rap', 'girl group', 'australian psych', 'pop dance', 'oulu metal', 'brill building pop', 'trap queen', 'rap cearense', 'eurovision', 'vocal jazz', 'electro house', 'piseiro', 'disco', 'trap chileno', 'neo-psychedelic', 'boy band', 'new wave', 'weirdcore', 'tropical house', 'gangster rap', 'modern country rock', 'chicago drill', 'emo rap', 'seattle hip hop', 'r&b brasileiro', 'pop venezolano', 'folk rock', 'argentine hip hop', 'album rock', 'rap conscient', 'trap italiana', 'trap soul', 'brooklyn drill', 'banda', 'grunge', 'australian dance', 'dance rock', 'metropopolis', 'k-pop boy group', 'singer-songwriter', 'german drill', 'rap marseille', 'belgian hip hop', 'funk bh', 'christlicher rap', 'nyc rap', 'latin pop', 'escape room', 'new romantic', 'celtic punk', 'house', 'celtic rock', 'folk punk', 'sertanejo pop', 'eurodance', 'bedroom soul', 'urbano espanol', 'pop', 'sheffield indie', 'chicago indie', 'madchester', 'gauze pop', 'australian hip hop', 'rock', 'german trance', 'funk']

Number of unique types of genres: 334

Now, we would like to export this list into a .csv file and have each element in the list be written row-wise.

Then our team will use the .csv file to map each of the genres to one of the 4 main genres.

```
[11]: # Importing library
import csv

# opening the csv file in 'w+' mode
file = open('genres.csv', 'w+', newline='')

# writing the data into the file
with file:
    write = csv.writer(file)
    write.writerows([unique_genres])
```

```
[12]: count = pd.Series(combined).value_counts()
print(count)
```

```
type(count)
```

```
pop          528
rap          311
dance pop    257
post-teen pop 256
latin        172
...
icelandic pop      1
sunnlensk tonlist  1
bubblegrunge       1
chicago indie     1
turkish trap       1
Length: 334, dtype: int64
```

```
[12]: pandas.core.series.Series
```

```
[13]: count.to_csv('genre_count.csv')
```

Our team has manually coded the unique genres above one of the 4 following genres: **pop**, **rap**, **rock** and **r&b**.

Now, we want to read the csv file into Python and create a dictionary based on the coded list so that we can map onto our Spotify dataset.

Note that since the DType of our **Genre** column was **object** and not **list**, we will first have to convert them into list before we conduct our mapping.

```
[14]: spotify['Genre'] = delister(spotify['Genre'])
spotify
```

```
[14]:
```

	Index	Highest Charting Position	Number of Times Charted	\
0	1	1	8	
1	2	2	3	
2	3	1	11	
3	4	3	5	
4	5	5	1	
...	
1540	1552	195	1	
1541	1553	196	1	
1542	1554	197	1	
1543	1555	198	1	
1544	1556	199	1	

	Week of Highest Charting	Song Name	Streams	\
0	2021-07-23--2021-07-30	Beggin'	48633449	
1	2021-07-23--2021-07-30	STAY (with Justin Bieber)	47248719	
2	2021-06-25--2021-07-02	good 4 u	40162559	

3	2021-07-02--2021-07-09	Bad Habits	37799456
4	2021-07-23--2021-07-30	INDUSTRY BABY (feat. Jack Harlow)	33948454
...
1540	2019-12-27--2020-01-03	New Rules	4630675
1541	2019-12-27--2020-01-03	Cheirosa - Ao Vivo	4623030
1542	2019-12-27--2020-01-03	Havana (feat. Young Thug)	4620876
1543	2019-12-27--2020-01-03	Surtada - Remix Brega Funk	4607385
1544	2019-12-27--2020-01-03	Lover (Remix) [feat. Shawn Mendes]	4595450

	Artist	Artist Followers	Song ID	\
0	Måneskin	3377762.0	3Wrjm47oTz2sjIgck1115e	
1	The Kid LAROI	2230022.0	5HCyWlXZPP0y6Gqq8TgA20	
2	Olivia Rodrigo	6266514.0	4ZtFanR9U6ndgddUvNcjcG	
3	Ed Sheeran	83293380.0	6PQ88X9TkUIAUIZJHW2upE	
4	Lil Nas X	5473565.0	27NovPIUIRrOZOCHxABJwK	
...	
1540	Dua Lipa	27167675.0	2ekn2ttSfGqwhhate0LSR0	
1541	Jorge & Mateus	15019109.0	2PWjKmJyTZedpmOUa3a5da	
1542	Camila Cabello	22698747.0	1rf0faqEpACxVEHIZBJe6W	
1543	Dadá Boladão, Tati Zaqui, OIK	208630.0	5F8ffc8KWNawllr5WsW0r	
1544	Taylor Swift	42227614.0	3i9UVldZOE0aD0JnyfAZZ0	

	Genre	...	Loudness	Speechiness	\
0	[italian pop, indie rock italiano]	...	-4.808	0.0504	
1	[australian hip hop]	...	-5.484	0.0483	
2	[pop]	...	-5.044	0.1540	
3	[pop, uk pop]	...	-3.712	0.0348	
4	[lgbtq+ hip hop, pop rap]	...	-7.409	0.0615	
...	
1540	[pop, uk pop, dance pop]	...	-6.021	0.0694	
1541	[sertanejo, sertanejo universitario]	...	-3.123	0.0851	
1542	[pop, electropop, dance pop, post-teen pop]	...	-4.333	0.0300	
1543	[funk carioca, brega funk]	...	-7.026	0.0587	
1544	[pop, post-teen pop]	...	-7.176	0.0640	

	Acousticness	Liveness	Tempo	Duration (ms)	Valence	Chord	\
0	0.12700	0.3590	134.002	211560.0	0.589	B	
1	0.03830	0.1030	169.928	141806.0	0.478	C#/Db	
2	0.33500	0.0849	166.928	178147.0	0.688	A	
3	0.04690	0.3640	126.026	231041.0	0.591	B	
4	0.02030	0.0501	149.995	212000.0	0.894	D#/Eb	
...	
1540	0.00261	0.1530	116.073	209320.0	0.608	A	
1541	0.24000	0.3330	152.370	181930.0	0.714	B	
1542	0.18400	0.1320	104.988	217307.0	0.394	D	
1543	0.24900	0.1820	154.064	152784.0	0.881	F	
1544	0.43300	0.0862	205.272	221307.0	0.422	G	

	Week of Highest Charting_Start	Week of Highest Charting_End
0	2021-07-23	2021-07-30
1	2021-07-23	2021-07-30
2	2021-06-25	2021-07-02
3	2021-07-02	2021-07-09
4	2021-07-23	2021-07-30
...
1540	2019-12-27	2020-01-03
1541	2019-12-27	2020-01-03
1542	2019-12-27	2020-01-03
1543	2019-12-27	2020-01-03
1544	2019-12-27	2020-01-03

[1545 rows x 25 columns]

```
[15]: # Creating our mapping dictionary
recode = {}
url_gcs = "https://raw.githubusercontent.com/jia-shing/pic16a-spotify/main/genre_count_sorted.csv" # Make sure the url is the raw version of the file on GitHub
gcs = pd.read_csv(url_gcs, header = None)

for index, row in gcs.iterrows():
    key = row[0]
    recode[key] = [[row[1]]]

recode
```

```
[15]: {'pop': [['pop']],
      'rap': [['rap']],
      'dance pop': [['pop']],
      'post-teen pop': [['pop']],
      'pop rap': [['pop']],
      'trap': [['rap']],
      'reggaeton': [['rap']],
      'trap latino': [['rap']],
      'latin': [['latin']],
      'hip hop': [['rap']],
      'melodic rap': [['rap']],
      'canadian pop': [['pop']],
      'electropop': [['pop']],
      'pop dance': [['pop']],
      'atl hip hop': [['rap']],
      'german hip hop': [['rap']],
      'uk pop': [['pop']],
      'edm': [['pop']],
```

'chicago rap': [['pop']],
'k-pop': [['pop']],
'reggaeton colombiano': [['rap']],
'tropical house': [['pop']],
'brooklyn drill': [['rap']],
'philly rap': [['rap']],
'rock': [['rock']],
'canadian hip hop': [['rap']],
'k-pop boy group': [['pop']],
'north carolina hip hop': [['rap']],
'toronto rap': [['rap']],
'r&b': [['r&b']],
'southern hip hop': [['rap']],
'detroit hip hop': [['rap']],
'canadian contemporary r&b': [['r&b']],
'modern rock': [['rock']],
'puerto rican pop': [['pop']],
'reggaeton flow': [['rap']],
'alt z': [['rock']],
'viral pop': [['pop']],
'adult standards': [['rock']],
'k-pop girl group': [['pop']],
'ohio hip hop': [['rap']],
'conscious hip hop': [['rap']],
'colombian pop': [['pop']],
'lounge': [['pop']],
'latin hip hop': [['rap']],
'slap house': [['pop']],
'pop rock': [['rock']],
'funk carioca': [['rock']],
'latin pop': [['pop']],
'trap queen': [['rap']],
'boy band': [['pop']],
'alternative r&b': [['r&b']],
'atl trap': [['rap']],
'contemporary country': [['pop']],
'emo rap': [['rap']],
'classic rock': [['rock']],
'miami hip hop': [['rap']],
'easy listening': [['pop']],
'sertanejo universitario': [['pop']],
'permanent wave': [['rock']],
'australian pop': [['pop']],
'vocal jazz': [['pop']],
'hamburg hip hop': [['rap']],
'pittsburgh rap': [['rap']],
'mellow gold': [['rock']],

'italian pop': [['pop']],
 'uk hip hop': [['rap']],
 'trap argentino': [['rap']],
 'electro house': [['pop']],
 'sertanejo': [['latin']],
 'dfw rap': [['rap']],
 'soft rock': [['rock']],
 'pop urbaine': [['pop']],
 'french hip hop': [['rap']],
 'uk dance': [['pop']],
 'sertanejo pop': [['pop']],
 'indie cafe pop': [['pop']],
 'cali rap': [['rap']],
 'panamanian pop': [['pop']],
 'talent show': [['pop']],
 'trap boricua': [['rap']],
 'australian hip hop': [['rap']],
 'underground hip hop': [['rap']],
 'nyc rap': [['rap']],
 'viral rap': [['rap']],
 'argentine hip hop': [['rap']],
 'lgbtq+ hip hop': [['rap']],
 'album rock': [['rock']],
 'canadian trap': [['rap']],
 'german cloud rap': [['rap']],
 'pop edm': [['pop']],
 'forro': [['latin']],
 'urban contemporary': [['pop']],
 'german pop': [['pop']],
 'italian hip hop': [['rap']],
 'house': [['pop']],
 'indie pop': [['pop']],
 'indie poptimism': [['pop']],
 'francoton': [['rap']],
 'rap francais': [['rap']],
 'pop venezolano': [['pop']],
 'german drill': [['rap']],
 'progressive house': [['pop']],
 'german dance': [['pop']],
 'r&b en espanol': [['r&b']],
 'escape room': [['pop']],
 'houston rap': [['rap']],
 'sad rap': [['rap']],
 'pop nacional': [['pop']],
 'neo mellow': [['pop']],
 'beatlesque': [['rock']],
 'rap conscient': [['rap']],

'new wave pop': [['pop']],
 'soul': [['r&b']],
 'vapor trap': [['rap']],
 'hip pop': [['rap']],
 'neo-psychedelic': [['rock']],
 'piseiro': [['rap']],
 'australian psych': [['rock']],
 'art pop': [['pop']],
 'australian dance': [['pop']],
 'trap italiana': [['rap']],
 'social media pop': [['pop']],
 'etherpop': [['pop']],
 'german trap': [['rap']],
 'nz pop': [['pop']],
 'kentucky hip hop': [['rap']],
 'gangster rap': [['rap']],
 'memphis hip hop': [['rap']],
 'funk pop': [['pop']],
 'dmv rap': [['rap']],
 'electronic trap': [['rap']],
 'indie rock italiano': [['rock']],
 'glam rock': [['rock']],
 'dutch edm': [['pop']],
 'progressive electro house': [['pop']],
 'bedroom pop': [['pop']],
 'modern alternative rock': [['rock']],
 'country road': [['pop']],
 'brostep': [['pop']],
 'art rock': [['rock']],
 'chill r&b': [['r&b']],
 'rockabilly': [['rock']],
 'pop soul': [['pop']],
 'europop': [['pop']],
 'dominican pop': [['pop']],
 'minnesota hip hop': [['rap']],
 'modern indie pop': [['pop']],
 'german alternative rap': [['rap']],
 'eurovision': [['pop']],
 'pop house': [['pop']],
 'piano rock': [['rock']],
 'pop reggaeton': [['pop']],
 'deep house': [['pop']],
 'deep euro house': [['pop']],
 'florida rap': [['rap']],
 'rap latina': [['rap']],
 'pagode baiano': [['pop']],
 'trap soul': [['rap']],

'deep underground hip hop': [['rap']],
 'urbano espanol': [['rap']],
 'brill building pop': [['pop']],
 'trap triste': [['rap']],
 'metropopolis': [['pop']],
 'new wave': [['rock']],
 'jazz pop': [['pop']],
 'country': [['pop']],
 'queens hip hop': [['rap']],
 'funk ostentacao': [['rock']],
 'regional mexican': [['rock']],
 'singer-songwriter': [['pop']],
 'motown': [['rock']],
 'london rap': [['rap']],
 'frauenrap': [['rap']],
 'basshall': [['rap']],
 'yacht rock': [['rock']],
 'comic': [['pop']],
 'new romantic': [['pop']],
 'indie pop rap': [['rap']],
 'belgian hip hop': [['rap']],
 'synthpop': [['pop']],
 'heartland rock': [['rock']],
 'funk': [['rock']],
 'quiet storm': [['rock']],
 'rap belge': [['rap']],
 'oakland hip hop': [['rap']],
 'trance': [['pop']],
 'hard rock': [['rock']],
 'rock-and-roll': [['rock']],
 'mariachi': [['latin']],
 'nuevo regional mexicano': [['latin']],
 'melodic metalcore': [['rock']],
 'metalcore': [['rock']],
 'uk metalcore': [['rock']],
 'girl group': [['pop']],
 'ranchera': [['latin']],
 'disco': [['pop']],
 'swing': [['rap']],
 'afroswing': [['rap']],
 'hardcore hip hop': [['rap']],
 'pop r&b': [['r&b']],
 'folk rock': [['rock']],
 'rhode island rap': [['rap']],
 'british soul': [['r&b']],
 'dembow': [['pop']],
 'plugg': [['rap']],

'brega funk': [['rock']],
'sheffield indie': [['pop']],
'indietronica': [['pop']],
'albanian hip hop': [['rap']],
'madchester': [['rock']],
'new french touch': [['pop']],
'italian adult pop': [['pop']],
'dream smp': [['pop']],
'britpop': [['pop']],
'rap marseille': [['rap']],
'folk-pop': [['pop']],
'bedroom soul': [['r&b']],
'funk paulista': [['rock']],
'modern country rock': [['pop']],
'east coast hip hop': [['rap']],
'rap dominicano': [['rap']],
'neo soul': [['r&b']],
'garage rock': [['rock']],
'meme rap': [['rap']],
'venezuelan hip hop': [['pop']],
'r&b brasileiro': [['r&b']],
'grime': [['rap']],
'south african house': [['pop']],
'moombahton': [['pop']],
'uk funky': [['pop']],
'big room': [['pop']],
'country rock': [['rock']],
'southern soul': [['r&b']],
'new orleans rap': [['rap']],
'deep german hip hop': [['rap']],
'classic soul': [['r&b']],
'chicago soul': [['r&b']],
'gauze pop': [['pop']],
'uk alternative hip hop': [['rap']],
'alternative pop rock': [['rock']],
'canadian latin': [['pop']],
'rap cearense': [['rap']],
'trap brasileiro': [['rap']],
'a cappella': [['pop']],
'swedish pop': [['pop']],
'indie r&b': [['r&b']],
'jawaiian': [['pop']],
'jazz funk': [['pop']],
'west coast rap': [['rap']],
'acoustic pop': [['pop']],
'shiver pop': [['pop']],
'trap chileno': [['rap']],

'perreo': [['latin']],
 'pop argentino': [['pop']],
 'christlicher rap': [['rap']],
 'indie rockism': [['rock']],
 'german underground rap': [['rap']],
 'k-rap': [['rap']],
 'chicago drill': [['rap']],
 'australian rock': [['rock']],
 'drill': [['rap']],
 'electro latino': [['pop']],
 'banda': [['latin']],
 'corrido': [['latin']],
 'norteno': [['latin']],
 'champeta': [['latin']],
 'scandipop': [['pop']],
 'brazilian hip hop': [['pop']],
 'soundtrack': [['pop']],
 'pacific islands pop': [['pop']],
 'tennessee hip hop': [['rap']],
 'dutch pop': [['pop']],
 'weirdcore': [['rock']],
 'afrofuturism': [['rock']],
 'uk alternative pop': [['pop']],
 'sunnlensk tonlist': [['pop']],
 'icelandic pop': [['pop']],
 'chicago indie': [['rock']],
 'bubblegrunge': [['rock']],
 'boston hip hop': [['rap']],
 'trancecore': [['rock']],
 'oulu metal': [['rock']],
 'latin viral pop': [['pop']],
 'nouvelle chanson francaise': [['pop']],
 'nyc pop': [['pop']],
 'funk 150 bpm': [['rock']],
 'rap metal': [['rap']],
 'post-grunge': [['rock']],
 'nu metal': [['rock']],
 'alternative metal': [['rock']],
 'seattle hip hop': [['rap']],
 'grunge': [['rock']],
 'irish singer-songwriter': [['r&b']],
 'norwegian pop': [['pop']],
 'dance rock': [['rock']],
 'folktronica': [['rock']],
 'eau claire indie': [['rock']],
 'rave funk': [['rock']],
 'sudanese pop': [['pop']],

```

'png pop': [['popo']],
'melanesian pop': [['pop']],
'torch song': [['pop']],
'vancouver indie': [['rock']],
'show tunes': [['pop']],
'hollywood': [['pop']],
'eurodance': [['pop']],
'classic uk pop': [['pop']],
'punk': [['rock']],
'irish folk': [['r&b']],
'folk punk': [['rock']],
'celtic rock': [['rock']],
'celtic punk': [['rock']],
'celtic': [['rock']],
'american folk revival': [['pop']],
'musical advocacy': [['pop']],
'dreamo': [['rock']],
'indie surf': [['rock']],
'cumbia pop': [['pop']],
'tekk': [['rock']],
'cubaton': [['pop']],
'funk bh': [['rock']],
'surf punk': [['rock']],
'vegas indie': [['pop']],
'old school rap francais': [['rap']],
'italian pop rock': [['rock']],
'italian indie pop': [['pop']],
'indonesian pop': [['pop']],
'german techno': [['rock']],
'german trance': [['rock']],
'aussietronica': [['rock']],
'mexican hip hop': [['pop']],
'rap conciencia': [['rap']],
'turkish trap': [['rap']]

```

```

[16]: # Remove extra list enclosing
for key, value in recode.items():
    recode[key] = value[0][0]

#Remapping each of the genres to the new genre based on the recode dictionary
for x in spotify['Genre']:
    for y in range(len(x)):
        x[y] = recode[x[y]]
spotify

```

```

[16]:      Index  Highest Charting Position  Number of Times Charted  \
0          1                               1                      8

```

1	2	2	3
2	3	1	11
3	4	3	5
4	5	5	1
...
1540	1552	195	1
1541	1553	196	1
1542	1554	197	1
1543	1555	198	1
1544	1556	199	1

	Week of Highest Charting	Song Name	Streams \
0	2021-07-23--2021-07-30	Beggin'	48633449
1	2021-07-23--2021-07-30	STAY (with Justin Bieber)	47248719
2	2021-06-25--2021-07-02	good 4 u	40162559
3	2021-07-02--2021-07-09	Bad Habits	37799456
4	2021-07-23--2021-07-30	INDUSTRY BABY (feat. Jack Harlow)	33948454
...
1540	2019-12-27--2020-01-03	New Rules	4630675
1541	2019-12-27--2020-01-03	Cheirosa - Ao Vivo	4623030
1542	2019-12-27--2020-01-03	Havana (feat. Young Thug)	4620876
1543	2019-12-27--2020-01-03	Surtada - Remix Brega Funk	4607385
1544	2019-12-27--2020-01-03	Lover (Remix) [feat. Shawn Mendes]	4595450

	Artist	Artist Followers	Song ID \
0	Måneskin	3377762.0	3Wrjm47oTz2sjIgck1115e
1	The Kid LAROI	2230022.0	5HCyWlXZPP0y6Gqq8TgA20
2	Olivia Rodrigo	6266514.0	4ZtFanR9U6ndgddUvNcjG
3	Ed Sheeran	83293380.0	6PQ88X9TkUIAUIZJHW2upE
4	Lil Nas X	5473565.0	27NovPIUIRrOZoCHxABJwK
...
1540	Dua Lipa	27167675.0	2ekn2ttSfGqwhhate0LSR0
1541	Jorge & Mateus	15019109.0	2PWjKmJyTZedpmOUa3a5da
1542	Camila Cabello	22698747.0	1rfofaqEpACxVEHIZBJe6W
1543	Dadá Boladão, Tati Zaqui, OIK	208630.0	5F8ffc8KWKNawllr5WsW0r
1544	Taylor Swift	42227614.0	3i9UVldZOE0aD0JnyfAZZ0

	Genre	...	Loudness	Speechiness	Acousticness	Liveness \
0	[pop, rock]	...	-4.808	0.0504	0.12700	0.3590
1	[rap]	...	-5.484	0.0483	0.03830	0.1030
2	[pop]	...	-5.044	0.1540	0.33500	0.0849
3	[pop, pop]	...	-3.712	0.0348	0.04690	0.3640
4	[rap, pop]	...	-7.409	0.0615	0.02030	0.0501
...
1540	[pop, pop, pop]	...	-6.021	0.0694	0.00261	0.1530
1541	[latin, pop]	...	-3.123	0.0851	0.24000	0.3330
1542	[pop, pop, pop, pop]	...	-4.333	0.0300	0.18400	0.1320

1543	[rock, rock]	...	-7.026	0.0587	0.24900	0.1820
1544	[pop, pop]	...	-7.176	0.0640	0.43300	0.0862

	Tempo	Duration (ms)	Valence	Chord	Week of Highest Charting_Start \
0	134.002	211560.0	0.589	B	2021-07-23
1	169.928	141806.0	0.478	C#/Db	2021-07-23
2	166.928	178147.0	0.688	A	2021-06-25
3	126.026	231041.0	0.591	B	2021-07-02
4	149.995	212000.0	0.894	D#/Eb	2021-07-23
...
1540	116.073	209320.0	0.608	A	2019-12-27
1541	152.370	181930.0	0.714	B	2019-12-27
1542	104.988	217307.0	0.394	D	2019-12-27
1543	154.064	152784.0	0.881	F	2019-12-27
1544	205.272	221307.0	0.422	G	2019-12-27

	Week of Highest Charting_End
0	2021-07-30
1	2021-07-30
2	2021-07-02
3	2021-07-09
4	2021-07-30
...	...
1540	2020-01-03
1541	2020-01-03
1542	2020-01-03
1543	2020-01-03
1544	2020-01-03

[1545 rows x 25 columns]

```
[17]: # Remove duplicates within each song's list of genres
for i, x in enumerate(spotify['Genre']):
    unique_genres = []
    for genre in x:
        if genre not in unique_genres:
            unique_genres.append(genre)
    # Update the 'Genre' column with the unique genres for this song
    spotify.at[i, 'Genre'] = unique_genres
```

There were multiple genres given in the original data that we felt we could not immediately classify into one of the 4 genres because they instead fell under a more general banner of Latin music, which was one of the most popular genres, but one where we could not immediately categorize entire genres such as “latin” to be mapped to either pop,rap, etc.

Instead in this step we created a **key** for latin that will allow us to later filter through and drop the key if another genre already exists for that song, and if not, then we would be able to look at the context of song/artist and make a manual decision later on.

```
[18]: #remove additional category of latin if there is another genre
for i, x in enumerate(spotify['Genre']):
    if "latin" in x and len(x) > 1:
        spotify.at[i, 'Genre'].remove("latin")

url_lr = "https://raw.githubusercontent.com/jia-shing/pic16a-spotify/main/
↳latin_remap.csv"
lr = pd.read_csv(url_lr, header = None)

for i, row in lr.iterrows():
    for j, name in enumerate(spotify['Song Name']):
        if row[0] == name:
            spotify.at[j, "Genre"] = row[1]

[19]: #identify which songs are listed under multiple genres to make a manual csv
↳recoding dictionary
for i, x in enumerate(spotify['Genre']):
    if len(x)>1:
        print(spotify.at[i, "Genre"])
```

```
['pop', 'rock']
['rap', 'pop']
['rap', 'pop']
['pop', 'rap']
['pop', 'rock']
['r&b', 'pop']
['r&b', 'pop']
['rap', 'pop']
['rap', 'pop']
['rap', 'pop']
['rock', 'pop']
rock
['pop', 'rap']
['pop', 'rock']
['pop', 'rock']
['pop', 'rap']
['rock', 'pop']
['r&b', 'pop']
['pop', 'r&b']
r&b
['pop', 'rock']
['pop', 'rock']
['rap', 'pop']
['pop', 'rap']
['pop', 'rap']
['popo', 'pop']
['rap', 'pop']
```



```

['rap', 'pop']
pop
['pop', 'rock']
['pop', 'rock']
rap
['rock', 'pop']
['pop', 'rap']
['rap', 'pop']
['rap', 'pop']
['pop', 'rap']
pop
['pop', 'rap']
['r&b', 'pop']
['pop', 'rock']
['pop', 'rock']
['pop', 'rock']
['pop', 'rock']
rock
['pop', 'rock']
['pop', 'rock', 'r&b']
['rock', 'pop']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
rock
['pop', 'rock']
['rock', 'pop']
pop
rock
['rap', 'pop']
['pop', 'rap']
['pop', 'r&b']
['pop', 'rap']
['pop', 'rap']
['pop', 'r&b']
['pop', 'rap']
['r&b', 'pop']
['pop', 'r&b']
['pop', 'rock']
['pop', 'rap']
pop
['pop', 'rock']
['rock', 'pop']
['pop', 'rap']
['rock', 'pop']
['rap', 'pop']
['pop', 'rap']
['pop', 'rap', 'r&b']

```

```

['pop', 'rap']
['rock', 'pop']
['pop', 'rap', 'r&b']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['rap', 'pop']
['rock', 'pop']
['pop', 'rap']
['pop', 'rock']
['pop', 'rap']
['rock', 'rap', 'pop']
['rap', 'pop']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['rock', 'rap']
['pop', 'rap']
['pop', 'rap']
['pop', 'rock']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['rap', 'pop']
['pop', 'rock']
['rap', 'pop']
['pop', 'rock']
pop
['rap', 'pop']
['rap', 'pop']
['pop', 'rap']
['rock', 'pop']
['pop', 'rap']
['rap', 'pop']
['rock', 'pop']
['rap', 'pop']
rap
['pop', 'rock', 'r&b']
['pop', 'rap']
['rap', 'pop']
['rap', 'pop']
['pop', 'r&b']
['pop', 'rap']
['rap', 'pop']
pop
['pop', 'rap']
['rap', 'pop']
['rap', 'pop']

```

['rap', 'pop']
['pop', 'rap']
['pop', 'rap']
['pop', 'rock']
['r&b', 'pop']
['pop', 'rap']
['r&b', 'pop']
['pop', 'rock']
pop
['rock', 'pop']
['rap', 'pop']
['rap', 'pop']
pop
['rap', 'pop']
['pop', 'rock']
['pop', 'rock']
['pop', 'rock']
['pop', 'rap']
['rock', 'pop']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['r&b', 'pop']
['pop', 'rap']
['pop', 'rap']
['rap', 'pop']
['pop', 'rap']
['rap', 'pop']
['r&b', 'pop']
['pop', 'rock']
['pop', 'rap']
['rap', 'pop']
['pop', 'rock']
['rap', 'pop']
['pop', 'rock']
['pop', 'rock']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['pop', 'rock']
['pop', 'rap']
['pop', 'rock']
['pop', 'rap']
r&b
['rap', 'pop']
['pop', 'rap']
['pop', 'rap']
['r&b', 'pop']

['rock', 'pop']
['rap', 'pop']
['r&b', 'pop']
['r&b', 'pop']
['rock', 'pop']
['r&b', 'pop']
['pop', 'rock']
['pop', 'rap']
['pop', 'rap']
['rap', 'rock']
['pop', 'rap', 'r&b']
['pop', 'rap']
['pop', 'rap']
['pop', 'r&b', 'rap']
['rap', 'pop']
['pop', 'rap']
['rap', 'pop']
['pop', 'r&b']
['rap', 'pop']
['pop', 'rap']
['pop', 'r&b']
['rock', 'pop']
['rock', 'pop']
['rock', 'pop']
['rock', 'pop']
['pop', 'rock']
['rock', 'pop']
['rock', 'pop']
['rock', 'pop']
['rock', 'pop', 'r&b']
['rap', 'rock', 'pop']
['rock', 'pop']
['rock', 'r&b']
['rock', 'pop']
['rock', 'pop']
['rap', 'pop']
['rock', 'pop']
['pop', 'rap']
['rock', 'pop']
['r&b', 'rock']
['rock', 'pop']
['rock', 'pop']
['rock', 'pop']
['pop', 'r&b']
['pop', 'rap']
['rock', 'pop']
['rock', 'pop']
['rock', 'pop']

```

['pop', 'r&b']
['rap', 'rock', 'pop']
['r&b', 'rock', 'rrock']
['rock', 'pop']
['pop', 'rap']
['rap', 'pop']
['rap', 'pop']
['rap', 'pop']
['rap', 'pop']
['rap', 'pop']
['rap', 'pop']
['rap', 'pop']
['rap', 'pop']
['rap', 'pop']
['rap', 'pop']
['rap', 'pop']
['r&b', 'pop']
['rap', 'r&b', 'pop']
['pop', 'rock']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['rap', 'pop']
['pop', 'rap']
['rock', 'pop']
['rap', 'pop']
['pop', 'rap']
['pop', 'rap']
['pop', 'r&b', 'rap']
['pop', 'rap']
['pop', 'rap']
['r&b', 'pop']
['rap', 'pop']
['pop', 'rap']
['pop', 'rap']
rap
['pop', 'r&b']
['pop', 'rock']
['rap', 'pop']
['pop', 'rap']
['rap', 'pop']
['pop', 'rock']
['pop', 'rock']
['pop', 'r&b']
['pop', 'rap']
['rap', 'pop']
['pop', 'rap']
['pop', 'rap']

```

```

['pop', 'rap']
['pop', 'rock']
['pop', 'rock']
['rap', 'pop']
['pop', 'rap']
['pop', 'r&b', 'rap']
['pop', 'r&b', 'rap']
['pop', 'rap']
['r&b', 'pop']
['r&b', 'pop']
['r&b', 'pop']
['rap', 'pop']
['r&b', 'pop']
['r&b', 'pop']
['r&b', 'pop']
['pop', 'rock']
['pop', 'rap']
['r&b', 'pop']
['r&b', 'pop']
['rap', 'pop']
['rap', 'pop']
['rap', 'r&b', 'pop']
['pop', 'rap']
['pop', 'rock', 'rrock', 'r&b']
['pop', 'rap']
['pop', 'rap']
['rap', 'pop']
['rap', 'pop']
['pop', 'rap']
['pop', 'rap']
['pop', 'r&b']
['pop', 'rock']
['rap', 'r&b']
['pop', 'r&b']
['rap', 'pop']
['pop', 'rap']
['pop', 'rap']
['rock', 'pop']
['r&b', 'pop']
['rap', 'pop']
['rap', 'pop']
['rap', 'pop']
['pop', 'rap']
['rap', 'pop']
pop
['pop', 'rap']
['pop', 'rap']

```

['pop', 'rock']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['rap', 'pop']
['pop', 'rap']
['rap', 'pop']
['rap', 'pop']
['pop', 'rap']
['pop', 'r&b', 'rap']
['rap', 'pop']
['rap', 'pop']
['rap', 'pop']
['rap', 'pop']
['pop', 'rap']
['rap', 'pop']
['rap', 'pop']
['rap', 'pop']
['rap', 'pop']
['rap', 'pop']
['r&b', 'pop', 'rap']
['rap', 'pop']
['rap', 'pop']
['rap', 'pop']
['rap', 'r&b']
['rap', 'pop']
['rap', 'pop']
['rap', 'pop']
['rap', 'pop']
['pop', 'rock']
['rap', 'pop']
['pop', 'rap']
['rock', 'pop']
['pop', 'rap']
['pop', 'r&b']
['pop', 'rap']
['pop', 'rap']
['pop', 'rock']
['pop', 'rap']
['pop', 'rap']
['rap', 'pop']
['pop', 'rock']
['pop', 'rock']
['rap', 'pop']
['pop', 'rock']
['pop', 'r&b', 'rap']
['pop', 'rap']

['pop', 'rap']
['pop', 'rap']
['r&b', 'pop']
['r&b', 'pop']
['pop', 'rap']
['rap', 'pop']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['rap', 'pop']
['pop', 'rap']
['pop', 'rock']
['pop', 'r&b']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['rap', 'pop']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['rap', 'r&b', 'pop']
['pop', 'rap']
['rap', 'pop']
['pop', 'rap']
['pop', 'rap']
['rock', 'pop']
['pop', 'rap']
['rap', 'pop']
['pop', 'rock']
['rap', 'r&b', 'pop']
['rap', 'r&b', 'pop']
['rap', 'pop']
['pop', 'rap']
['r&b', 'pop']
['r&b', 'pop']
['r&b', 'pop']
['r&b', 'pop']
['r&b', 'pop']
['r&b', 'pop']
['r&b', 'pop']


```

['r&b', 'pop']
['r&b', 'pop']
['r&b', 'pop']
['pop', 'r&b']
['pop', 'rap']
['pop', 'r&b']
['pop', 'rap']
['rap', 'pop']
['pop', 'rap']
['rap', 'pop']
['rap', 'r&b']
['pop', 'rap']
['rap', 'pop']
['pop', 'rap']
['rap', 'r&b', 'pop']
['rap', 'pop']
pop
['rap', 'r&b']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['pop', 'r&b']
['pop', 'rap']
['pop', 'rap']
['rap', 'pop']
['pop', 'rap']
pop
['rap', 'pop']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['rap', 'pop']
['rap', 'pop']
['pop', 'rap']
['rap', 'pop']
['rap', 'pop']
['rap', 'pop']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['pop', 'rap']
['pop', 'rock']
['rap', 'pop']
['rap', 'pop']
['rap', 'pop']

```

In this step, we had to manually look through each song that was listed under multiple genres, due to both the original data set including multiple genres, so even after our encoding, it was likely

that a song fell under multiple categories. The approach to doing this was tedious but involved listening to the song itself in addition to looking at the artist and the original hyper-specific genres which sometimes gave more clues as to how a song would be categorized

Due to the nature of this task, it was highly subjective, as there was no function or program we could write to pick if a song would be qualified as “Rap” or “R&B”. There is a chance that if an individual were to conduct this categorization again that it could influence the way our machine learning model would predict outcomes.

One method we tried to employ during this process is having the same individual focused on the encoding of genres, so that it stayed relatively consistent throughout the project, rather than varying person by person.

```
[20]: url_sg = "https://raw.githubusercontent.com/jia-shing/pic16a-spotify/main/
↳single_genre.csv"
sg= pd.read_csv(url_sg, header = None)

for i, row in sg.iterrows():
    for j, name in enumerate(spotify['Song Name']):
        if row[0] == name:
            spotify.at[j, 'Genre']= row[3]
for i, val in enumerate(spotify['Genre']):
    if isinstance(val, list):
        spotify.at[i, 'Genre']=val[0]
```

After looking more closely at the data, we found a discrepency in songs that were released on 8/10/21-8/13/21: their popularity scores only ranged from 0-3 with one that was 15. This is very abnormal when 1), compared to the rest of the data where the lowest popularity score on the entire list was 76, and 2), because the other statistics such as highest charting position, number of times charted, streams, etc were much much higher than average when compared to other songs with lower popularity scores, which is why we believe that this is a glitch. Since we know that popularity could be a factor that we may use later, we wanted to remove these few outliers (total of 45).

Included an example of why we filtered these out below:

```
[21]: #although 8/13 is just one of the days, you can see examples of songs that by
↳other metrics are very "popular", but strangely have a score of 0
mask = spotify['Release Date'] == '8/13/21'
#print(spotify.loc[mask])

#one obvious example is the song "bad guy" by Billie Eilish
mask2 = spotify['Song Name'] == 'bad guy'
print(spotify.loc[mask2])

#When looking at the data, this song was charted 83 times, making it the 3rd
↳highest song to be
#charted out of the entire dataset, yet it has a popularity val of 0,
↳indicating some error in how
```

```
#the data was collected from Spotify's end
```

```
Index Highest Charting Position Number of Times Charted \
162    165                      13                      83

Week of Highest Charting Song Name Streams Artist \
162  2020-01-24--2020-01-31 bad guy 5436286 Billie Eilish

Artist Followers Song ID Genre ... Loudness Speechiness \
162    1250353.0 1hewNsVmijBqjKvFRQfk4m pop ... -10.965 0.375

Acousticness Liveness Tempo Duration (ms) Valence Chord \
162    0.328    0.1 135.128 194088.0 0.562 G

Week of Highest Charting_Start Week of Highest Charting_End
162                2020-01-24                2020-01-31
```

```
[1 rows x 25 columns]
```

```
[22]: #Drop songs
mask = spotify['Popularity'] < 16
num_rows_dropped = len(spotify[mask])
spotify.drop(spotify[mask].index, inplace=True)
spotify.drop(columns=['Index'], inplace=True)
print(num_rows_dropped)
```

```
45
```

4 Part 3: Exploratory Analysis

In this section, we want to compute some summary statistics and construct visualizations about the relationships between variables.

We will create 4 figures (with one figure with multiple axes) and 1 display table. They are as follows: - A (linear) correlation matrix with all quantitative variables - A display table summarizing quantitative variables that were relatively linearly correlated from the correlation matrix - A histogram displaying how many songs in the Top 50 were charted on each day of the month - A histogram and KDE pairplot to test our hypothesis that more upbeat and danceable songs would be more popular - A line plot with multiple axes (based on Genre) to see how Streams and Popularity (the two quantitative variables that proved significant thus far in our EDA) are correlated.

```

```

We will also add a new column to the spotify dataframe called “If top 50” and populates it with values based on the “Highest Charting Position” column. So the end result is that the “If top 50” column will contain TRUE for each row where the “Highest Charting Position” is less than or equal to 50, and a FALSE otherwise.

```
[23]: import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import seaborn as sns
%config InlineBackend.figure_format = 'retina'

spotifydf = spotify.copy()
spotifydf['If top 50'] = spotifydf['Highest Charting Position'].apply(lambda x:
    ↪ True if x <= 50 else False)

spotifydf
```

```
[23]:
```

	Highest Charting Position	Number of Times Charted	\
0	1	8	
1	2	3	
2	1	11	
3	3	5	
4	5	1	
...	
1540	195	1	
1541	196	1	
1542	197	1	
1543	198	1	
1544	199	1	

	Week of Highest Charting	Song Name	Streams	\
0	2021-07-23--2021-07-30	Beggin'	48633449	
1	2021-07-23--2021-07-30	STAY (with Justin Bieber)	47248719	
2	2021-06-25--2021-07-02	good 4 u	40162559	
3	2021-07-02--2021-07-09	Bad Habits	37799456	
4	2021-07-23--2021-07-30	INDUSTRY BABY (feat. Jack Harlow)	33948454	
...	
1540	2019-12-27--2020-01-03	New Rules	4630675	
1541	2019-12-27--2020-01-03	Cheirosa - Ao Vivo	4623030	
1542	2019-12-27--2020-01-03	Havana (feat. Young Thug)	4620876	
1543	2019-12-27--2020-01-03	Surtada - Remix Brega Funk	4607385	
1544	2019-12-27--2020-01-03	Lover (Remix) [feat. Shawn Mendes]	4595450	

	Artist	Artist Followers	Song ID	\
0	Måneskin	3377762.0	3Wrjm47oTz2sjIgck1115e	
1	The Kid LAROI	2230022.0	5HCyWlXZPP0y6Gqq8TgA20	
2	Olivia Rodrigo	6266514.0	4ZtFanR9U6ndgddUvNcjCg	
3	Ed Sheeran	83293380.0	6PQ88X9TkUIAUIZJHW2upE	
4	Lil Nas X	5473565.0	27NovPIUIRrOZOCHxABJwK	
...	
1540	Dua Lipa	27167675.0	2ekn2ttSfGqwhhate0LSR0	

1541	Jorge & Mateus	15019109.0	2PWjKmJyTZeDpmOUa3a5da
1542	Camila Cabello	22698747.0	1rfofaqEpACxVEHIZBJe6W
1543	Dadá Boladão, Tati Zaqui, OIK	208630.0	5F8ffc8KWKNAwllr5WsW0r
1544	Taylor Swift	42227614.0	3i9UVldZOE0aD0JnyfAZZ0

	Genre	Release Date	...	Speechiness	Acousticness	Liveness	Tempo	\
0	rock	12/8/17	...	0.0504	0.12700	0.3590	134.002	
1	rap	7/9/21	...	0.0483	0.03830	0.1030	169.928	
2	pop	5/21/21	...	0.1540	0.33500	0.0849	166.928	
3	pop	6/25/21	...	0.0348	0.04690	0.3640	126.026	
4	rap	7/23/21	...	0.0615	0.02030	0.0501	149.995	
...	
1540	pop	6/2/17	...	0.0694	0.00261	0.1530	116.073	
1541	pop	10/11/19	...	0.0851	0.24000	0.3330	152.370	
1542	pop	1/12/18	...	0.0300	0.18400	0.1320	104.988	
1543	rock	9/25/19	...	0.0587	0.24900	0.1820	154.064	
1544	pop	11/13/19	...	0.0640	0.43300	0.0862	205.272	

	Duration (ms)	Valence	Chord	Week of Highest Charting_Start	\
0	211560.0	0.589	B	2021-07-23	
1	141806.0	0.478	C#/Db	2021-07-23	
2	178147.0	0.688	A	2021-06-25	
3	231041.0	0.591	B	2021-07-02	
4	212000.0	0.894	D#/Eb	2021-07-23	
...	
1540	209320.0	0.608	A	2019-12-27	
1541	181930.0	0.714	B	2019-12-27	
1542	217307.0	0.394	D	2019-12-27	
1543	152784.0	0.881	F	2019-12-27	
1544	221307.0	0.422	G	2019-12-27	

	Week of Highest Charting_End	If top 50
0	2021-07-30	True
1	2021-07-30	True
2	2021-07-02	True
3	2021-07-09	True
4	2021-07-30	True
...
1540	2020-01-03	False
1541	2020-01-03	False
1542	2020-01-03	False
1543	2020-01-03	False
1544	2020-01-03	False

[1500 rows x 25 columns]


4.0.1 Visualization 1: Correlation Matrix

For our first visualization, we will plot a correlation matrix to see if there are any **linear correlation** that can be captured since those are the most direct relationship to establish between quantitative variables

```
[24]: # Computing the correlation for the train set
corr = spotifydf.corr()

# Plotting the correlation matrix
plt.figure(figsize = (15, 15))
sns.heatmap(corr, annot=True, cmap='coolwarm')
```

[24]: <AxesSubplot:>



Analysis

From the matrix, there doesn't seem to be any very **strong** linear correlations with `If top 50`, but there are specific variables that stand out that show some level of correlation. `Number of Times Charted`, `Streams`, `Artist Followers` and `Popularity` showed higher correlation coefficient values.

This chart also shows us that there are certain variables that seem to be relatively **linearly uncorrelated**; these were typically qualities of the song itself, like a song's `Acousticness` or `Duration`. This begins to explain that a song making it to the top 50 is not strongly related to inherent aspects of the song

Note that `Highest Charting Position` counter-intuitively has a strong negative correlation because `If top 50` is a boolean variable that is dependent on `Highest Charting Position`.

4.0.2 Visualization 2: Display Tables

In this section, we will select the variables that have a significant correlation with the `If top 50` variable and summarize the mean, standard deviation as well as size.

```
[25]: spotify_filtered = spotifydf[['Genre', 'If top 50', 'Streams', 'Artist_
↳Followers', 'Number of Times Charted', 'Popularity']]
summary = spotify_filtered.groupby(["Genre", "If top 50"]).aggregate([np.mean,
↳np.std, np.size])

summary
```

```
[25]:
```

			Streams			Artist Followers \	
			mean	std size		mean	
Genre	If top 50						
pop	False	5.805389e+06	1.451026e+06	428		1.549960e+07	
	True	7.768838e+06	5.378253e+06	247		2.398852e+07	
r&b	False	5.824660e+06	1.413711e+06	37		1.071472e+07	
	True	6.167265e+06	1.637651e+06	13		2.067369e+07	
rap	False	5.856299e+06	1.466847e+06	477		1.162166e+07	
	True	6.949243e+06	5.007990e+06	201		1.635432e+07	
rock	False	5.717884e+06	1.251989e+06	81		7.037229e+06	
	True	1.200093e+07	1.113141e+07	16		3.522286e+06	

			Number of Times Charted			\	
			std size	mean		std size	
Genre	If top 50						
pop	False	1.699299e+07	428	6.759346	11.033928	428	
	True	2.092276e+07	247	21.526316	22.730981	247	
r&b	False	1.433124e+07	37	5.621622	13.200066	37	
	True	1.656379e+07	13	10.307692	9.927946	13	
rap	False	1.372974e+07	477	4.794549	8.719950	477	

	True	1.541202e+07	201	19.358209	20.009274	201
rock	False	9.649465e+06	81	10.370370	15.874070	81
	True	5.150205e+06	16	11.375000	13.700973	16

		Popularity		std	size
		mean			
Genre	If top 50				
pop	False	70.899533	10.634956	428	
	True	77.080972	9.898752	247	
r&b	False	71.405405	8.411168	37	
	True	74.461538	8.875203	13	
rap	False	68.660377	8.890437	477	
	True	77.492537	7.459973	201	
rock	False	72.493827	13.325843	81	
	True	66.750000	18.142032	16	

Analysis

We wanted to use this table to explore what some of the influence of genre is on the following variables: **Streams**, **Artist Followers**, **# Times Charted** and **Popularity**; the variables we saw above to be linearly correlated with a song in the top 50

From the summary table, we can first observe that for all genres **except Rock**, that the mean streams, followers, times charted, and popularity were **always** higher if the song charted in the top 50 vs when it did not. Looking at **Rock** more closely allows for interesting insights into overall song trends: The mean number of artist followers for rock music that is not in the top 50 is only about half that of pop music that is not in the top 50. This suggests that rock music may have a smaller fan base or may be less popular among streaming music audiences. Additionally, rock music that is not in the top 50 has the lowest mean popularity score out of all the genres, suggesting that it may be harder for rock music to gain mainstream popularity compared to other genres. This is particularly evident when comparing it to pop music, which consistently has the highest mean popularity scores across all categories. Because of these factors, it indicates that it may be more challenging for rock music to gain mainstream popularity and chart on popular music charts.

Other notable observations include: - Pop music tends to have the largest number of followers, the most frequent charting, and the highest popularity scores, both for music that is in the top 50 and for music that is not in the top 50. - Rap music that is in the top 50 has the highest mean number of followers, while rap music that is not in the top 50 has the lowest mean number of followers. - Rock music that is not in the top 50 has the lowest mean popularity score, while rock music that is in the top 50 has the lowest mean number of followers. - R&B music has relatively consistent means for all three variables, regardless of whether the music is in the top 50 or not.

4.0.3 Visualization 3: Histogram

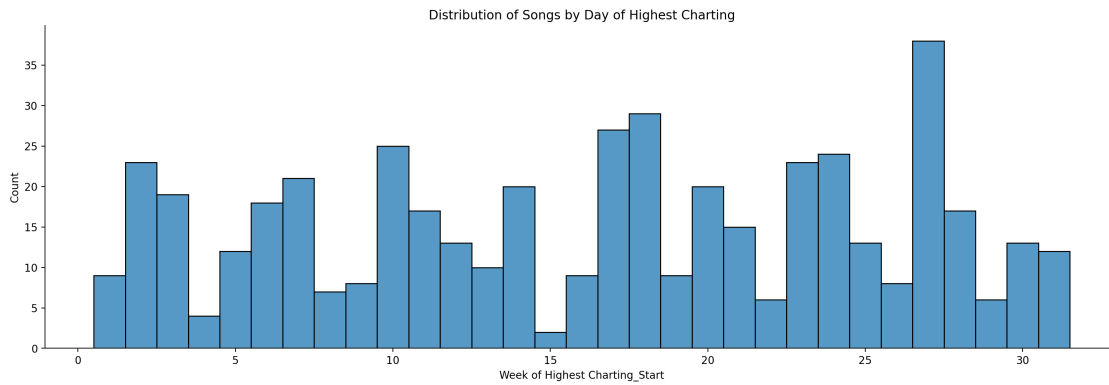
In this section, we want to see if day of the month affects charting position. Subset top 50 songs and see how many are charted on a particular day.

```
[26]: # This works, just run it again. Remove this after you read it.
spotify_top50 = spotifydf[spotifydf['If top 50'] == True]
```



```
year = spotify_top50['Week of Highest Charting_Start'].dt.day
sns.displot(year, discrete = True, aspect = 3, height = 5, kind = 'hist').
    ↪set(title = 'Distribution of Songs by Day of Highest Charting')
```

[26]: <seaborn.axisgrid.FacetGrid at 0x7fc22c600a60>



The histogram shows the distribution of top 50 songs based on the day of the month that they were charted. The data suggests that there is no clear pattern or relationship between the day of the month and the count of charted songs. Although there are certain days that appear to have higher counts, such as the 1st and the 15th, the data is too variably spread to conclude that a certain day of the month is important in determining a song's success.

This observation is supported by the fact that the histogram shows a relatively uniform distribution of counts across the days of the month, with no clear spikes or dips. This suggests that other factors, such as the song's genre, artist, and marketing strategy, may play a more important role in determining its success than the specific day of the month on which it was charted.

4.0.4 Visualization 4: Histogram and KDE

Since the Correlation Matrix did not provide us with expected results (e.g. more danceable and upbeat songs get streamed more and are charted higher), we suspect that they may not be linearly correlated but there is some other underlying relationship.

```
[27]: interested_features = ["Danceability", "Valence", "Tempo", "Energy",
    ↪ "Speechiness"]

fig, ax = plt.subplots(nrows=5, ncols=1, figsize=(20, 20))

for i in range(len(interested_features)):
    feature = interested_features[i]
    positive = spotifydf[spotifydf["If top 50"] == 1][feature]
    negative = spotifydf[spotifydf["If top 50"] == 0][feature]
```

```






sns.histplot(positive, bins=30, label="Positive", color="green", kde=True,
↪ax=ax[i], alpha=0.6)
sns.histplot(negative, bins=30, label="Negative", color="red", kde=True,
↪ax=ax[i], alpha=0.3)

ax[i].legend(loc='upper right')
ax[i].set_title(f"Histogram & KDE Plots for {feature}")

plt.subplots_adjust(left=0.1, right=0.9, top=0.95, bottom=0.05, hspace=0.5)

plt.show()

```



The histograms for energy, speechiness, danceability, and valence show that the positive and negative values look nearly identical in shape, with the main difference being scale. Which suggests that these traits are not **just** specific to successful songs that make it in the top 50. Rather, these traits are more general and make a song more palatable to a wider audience. For example, the histograms for energy, speechiness, and danceability show a symmetrical shape with one main peak, suggesting a normal distribution; this indicates that there are specific values of these traits that act as a sweet spot for the success of a song, as a songs with 100% energy are not likely popular. Nevertheless, the width of the curves suggests that the data is moderately spread out, meaning that there is some variability in the values of these traits among popular songs.

The striking similarities between the positive and negative histograms suggest that these traits are not helpful in differentiating the success of a song, which further reinforces the low correlation values seen in the earlier matrix analysis.

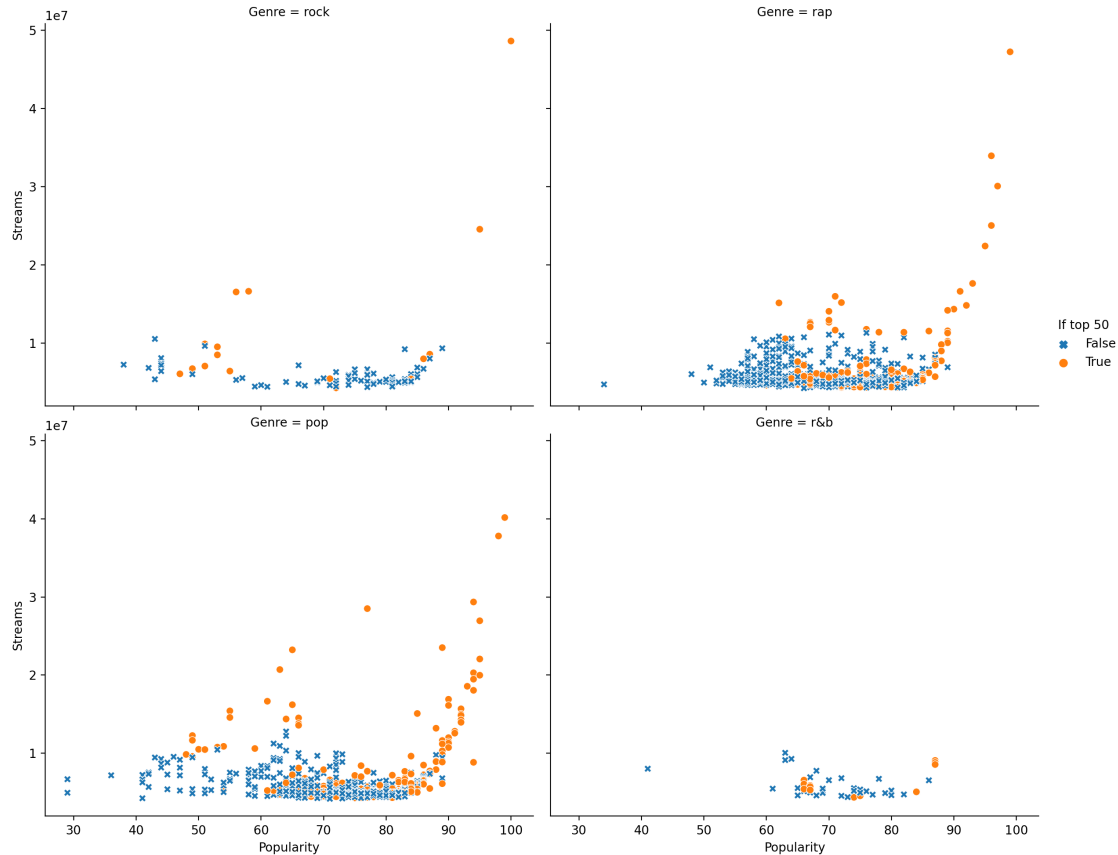
4.0.5 Visualization 5: Lineplot

Since Streams and Popularity both seemed to be variables that we positively correlated with the overall success of a song, we want to more closely look at how it compared between songs that made it into the top 50 vs those that didn't, and also split up the data by genre once again so we can visualize these differences in more specific context

```
[28]: fig3 = sns.relplot(data = spotifydf, x = "Popularity",
                        y = "Streams",
                        col = "Genre",
                        style = "If top 50",
                        hue = "If top 50",
                        kind = "scatter",
                        height = 6,
                        aspect = 1,
                        col_wrap = 2,
                        style_order = [True, False]
                        )
fig3.fig.subplots_adjust(top = 0.8)
fig3.fig.suptitle("Relationship between Popularity and Streams of Top 50 and
↳Non-Top 50 Songs by Genre on Spotify")
```

```
[28]: Text(0.5, 0.98, 'Relationship between Popularity and Streams of Top 50 and Non-
Top 50 Songs by Genre on Spotify')
```

Relationship between Popularity and Streams of Top 50 and Non-Top 50 Songs by Genre on Spotify



This collection of lineplots show us data that is mostly as we would have expected, but the separation into genres help to visualize the relationships we observed earlier in the `data` table to provide a more detailed breakdown of the correlations shown in the `correlation matrix`, but now adding in the additional key of genres.

Both `pop` and `rap` show very similar overall trends in how the relationship between popularity and streams is almost exponential, and almost exclusively followed by songs that made it in the top 50. While this seems intuitive, it is an important metric to observe because we do not necessarily know how the score for popularity is collected/calculated, so visualizing how strongly it is related to streams and if a song made it into the top 50 is helpful for us moving forward that it can be a variable we utilize.

The categories of `rock` and `R&B` do not provide such a clear pattern. Both categories never reached near the maximal scores of popularity and we can visually see that number of songs in these genres is far less than the size of the other two. This makes the analysis of certain trends within these categories slightly more difficult because we are working with a much smaller sample size. In addition, the proportion of songs from these genres that made it into the top 50 is not as large as

the pop and rap categories, so the spread of data from rock and R&B make it more difficult to identify outliers or what the general trend between popularity and streams are

5 Part 4: Feature Selection Preparation

In this section, we will use one qualitative feature and two quantitative features to choose the best data columns for our models. We will justify our choices and explain our approach, which we will repeat for each of the three models that we will build. We will use automated feature selection to estimate which combinations of three features work best for our models in the next section.

In order to do this, we had to drop the following columns from our dataframe: “Artist”, “Song Name”, “Release Date”, “Song ID”, “Weeks Charted”, “Week of Highest Charting_Start”, “Week of Highest Charting_End”, “Week of Highest Charting”, “Highest Charting Position”.

These are categories we decided are not relevant in deciding the success of a song because categories like Artist, Song Name, Song ID are aspects that will likely distract the success of our model. Artist name is not relevant since we know how famous they are from **Artist Followers**. We don’t want the model to learn whether someone with the first name “Ariana” or “Bruno” will definitely come up with popular songs since we have better measures. Song name and Song ID are both not important based on the same reasoning above.

The metric for Weeks charted we decided was too similar and not as clear or important a statistic as **Number of Times Charted**, so we dropped that as well.

In **Visualization 3**, we showed that the variable Week of highest Charting is not important, so we decided on not using that as a variable in our model.

The Highest Charting Position was dropped because we replaced it with the **If top 50** boolean column

```
[29]: #Delete columns that are not related to the characteristics of the song.
spotify_selection = spotifydf.drop(["Artist",
                                   "Song Name",
                                   "Release Date",
                                   "Song ID",
                                   "Weeks Charted",
                                   "Week of Highest Charting_Start",
                                   "Week of Highest Charting_End",
                                   "Week of Highest Charting",
                                   "Highest Charting Position"], axis=1)

#display data set
spotify_selection.head()
```

```
[29]:   Number of Times Charted   Streams  Artist Followers  Genre  Popularity  \
0                8  48633449        3377762.0  rock      100.0
1                3  47248719        2230022.0  rap       99.0
2               11  40162559        6266514.0  pop       99.0
3                5  37799456        83293380.0  pop       98.0
4                1  33948454        5473565.0  rap       96.0
```

	Danceability	Energy	Loudness	Speechiness	Acousticness	Liveness	\
0	0.714	0.800	-4.808	0.0504	0.1270	0.3590	
1	0.591	0.764	-5.484	0.0483	0.0383	0.1030	
2	0.563	0.664	-5.044	0.1540	0.3350	0.0849	
3	0.808	0.897	-3.712	0.0348	0.0469	0.3640	
4	0.736	0.704	-7.409	0.0615	0.0203	0.0501	

	Tempo	Duration (ms)	Valence	Chord	If top 50
0	134.002	211560.0	0.589	B	True
1	169.928	141806.0	0.478	C#/Db	True
2	166.928	178147.0	0.688	A	True
3	126.026	231041.0	0.591	B	True
4	149.995	212000.0	0.894	D#/Eb	True

To prepare the data set prior to inputting it into our machine learning models, we have to split it into training and test sets and clean the data. This is to ensure that the information from the data cleaning process does not unintentionally pollute the test set.

To split the data set into training and test sets, we used the `train_test_split()` function with a 20/80 split, such that 80% of the data set will be used for training and 20% will be used for testing. The training set will be used to develop a trained model, while the test set will be used to determine the accuracy of the model in predicting whether a song is in the Top 50.

```
[30]: from sklearn.model_selection import train_test_split

# Splitting the data into test set and train set.
# Using proportion of 80% of the data as train set, and 20% of the data as test
    ↪ set.
train, test = train_test_split(spotify_selection, test_size = 0.2, random_state
    ↪ = 20)

# Display the shape of train and test sets.
train.shape, test.shape
```

```
[30]: ((1200, 16), (300, 16))
```

In the code below, We tranfrom **Genre** and **Chord** into number, so that we can use it in feature selection and modeling.

Corresponding reference table for **Genre** - 0 = pop - 1 = r&b - 2 = rap - 3 = rock

Corresponding reference table for **Chord** - 0 = A - 1 = A#/Bb - 2 = B - 3 = C - 4 = C#/Db - 5 = D - 6 = D#/Eb - 7 = E - 8 = F - 9 = F#/Gb - 10 = G - 11 = G#/Ab

Here, we want to clean the dataset by integer encoding the qualitative variables (**Chord** and **Genre**) and split the data into **X** and **y** which contains the predictor and target variable(s) dataframe respectively.

We will write a function called `spotify_prep` to avoid repeating code.

```
[31]: from sklearn.preprocessing import LabelEncoder

def spotify_prep(data, target):
    """Prepare the data for modeling.
    Step 1: Make a copy of the data
    Step 2: Turn genre and chord into number
    Step 3: Split the data into X and y

    Args:
        data: train or test set
        target: target variable

    Returns:
        X: Predictor dataframe - cleaned df without target variable (i.e. If top 50)
        y: Target variable - If top 50 column only from cleaned df
    """

    # Makes a copy of data
    df = data.copy()

    # Turning genre and chord into number
    le = LabelEncoder()
    df['Genre'] = le.fit_transform(df['Genre'])
    df['Chord'] = le.fit_transform(df['Chord'])

    # Split the data into X and y
    X = df.drop([target], axis = 1)
    y = df[target]
    return (X,y)

# Getting train and test sets ready for modeling
X_train, y_train = spotify_prep(train, target = "If top 50")
X_test, y_test = spotify_prep(test, target = "If top 50")

# Display the shape of train and test sets.
print(X_train.shape, y_train.shape)
print(X_test.shape, y_test.shape)
```

```
(1200, 15) (1200,)
(300, 15) (300,)
```

```
[32]: # import necessary functions from sklearn
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import cross_val_score
```

```

from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import StratifiedKFold
from sklearn.svm import SVC

```

```

[33]: # ignore warnings
import warnings
from sklearn.exceptions import ConvergenceWarning
warnings.filterwarnings("ignore", category=ConvergenceWarning)
from sklearn.exceptions import FitFailedWarning
warnings.filterwarnings('ignore', category=FitFailedWarning)
warnings.filterwarnings("ignore", category=UserWarning)
from warnings import simplefilter
warnings.simplefilter(action='ignore', category=FutureWarning)

```

```

[34]: # define "spotify_quantitative" for all possible quantitative values
spotify_quantitative = spotify_selection.drop(["Genre", "Chord", "If top 50"],
axis=1)
spotify_qualitative = spotify_selection[["Genre", "Chord"]]
all_qualitative = list(spotify_qualitative.columns)
all_quantitative = list(spotify_quantitative.columns)

```

```

[35]: all_qualitative

```

```

[35]: ['Genre', 'Chord']

```

```

[36]: all_quantitative

```

```

[36]: ['Number of Times Charted',
'Streams',
'Artist Followers',
'Popularity',
'Danceability',
'Energy',
'Loudness',
'Speechiness',
'Acousticness',
'Liveness',
'Tempo',
'Duration (ms)',
'Valence']

```

Right now, we have clean training data ready to train with different models. In order to determine which model can highly predict if a song is in top 50 before, we must select the best three features for each model. For the qualitative feature, we can only use **Genre** and **chord**, as it is the only qualitative datas in our dataset. For the two quantitative features, we will loops through all available quantitative values to determine which two of them work best with **Genre** and **chord**.

We are going to create some functions that are useful for our features section. They are: -


```
get_cv_score(model) - display_best_cols(results)
```

```
[37]: def get_cv_score(model):  
    """  
    Calculate the scores for models using different pairs of quantitative_  
    ↪ features.  
  
    Args:  
        model: a function that takes a list of column names as input and  
        returns the score of the model using the given columns.  
  
    Returns:  
        A dictionary containing the scores for models using different pairs of_  
    ↪ quantitative features.  
    """  
    # create a dictionary to store the results  
    results = {}  
    # loop through all possible pairs of quantitative features  
    for i in range(len(all_quantitative)):  
        for j in range(i + 1, len(all_quantitative)):  
            for k in all_qualitative:  
                cols = [k, all_quantitative[i], all_quantitative[j]]  
                # score = np.round(model(cols), 3)  
                score = np.round(model(cols), 3)  
                results[str(cols)] = score  
                # print the results for the current set of columns  
                print(f"Training with columns {cols} \n The cv score is_  
    ↪ {score} \n")  
                print(type(score))  
    return results
```

```
[38]: def display_best_cols(results):  
    """  
    This function takes in the results dictionary of a model's scores and_  
    ↪ returns the combination of columns that resulted in the highest score along_  
    ↪ with the score itself.  
  
    Args:  
        results: A dictionary with column combinations as keys and their_  
    ↪ corresponding scores as values.  
  
    Returns:  
        None. The function simply prints out the combination with the highest_  
    ↪ score and its corresponding score.  
    """  
    # Get the highest cols combination.  
    best_cols = max(results, key=results.get)
```

```

best_score = results[best_cols]
print(f"The combination with the highest score is {best_cols} with a cv_
↳score of {best_score}.")

```

6 Part 5: Modelling

6.0.1 5.0: Plotting Decision Regions

```

[39]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

def plot_decision_regions(model, X_test, y_test, colx, coly,
↳qualitative_feature):
    """
    Plots the decision regions of a classifier along with the data points.

    Args:
    - model: a scikit-learn classifier that has a `fit` and `predict` method
    - X_test: a numpy array of shape (n_samples, n_features) representing the_
↳test data
    - Y_test: a numpy array of shape (n_samples,) representing the labels of_
↳the test data
    - colx: a string representing the column name for the x-axis of the plot
    - coly: a string representing the column name for the y-axis of the plot
    - qualitative_feature: a string representing the column name for the_
↳qualitative feature used to plot decision regions

    Returns:
    - None
    """
    # create dictionary for labels based on qualitative feature
    if qualitative_feature == "Genre":
        key = {0: "pop", 1: "r&b", 2: "rap", 3: "rock"}
        #iterate = X_test[qualitative_feature].max()
    elif qualitative_feature == "Chord":
        key = {0: "A", 1: "A#/Bb", 2: "B", 3: "C", 4: "C#/Db", 5: "D", 6: "D#/
↳Eb", 7: "E", 8: "F", 9: "F#/Gb", 10: "G", 11: "G#/Ab"}
        #iterate = X_test[qualitative_feature].max()+1
    # create dictionary for legend labels
    top_50_key = {0: "Not Top 50", 1: "Top 50"}

    for i in range(X_test[qualitative_feature].max()+1):
        # extract data for current qualitative feature value
        XX = X_test[[colx, coly]][X_test[qualitative_feature] == i]

```

```

index = X_test.index[X_test[qualitative_feature] == i].tolist()
YY = y_test[index]

# fit model to current data
model.fit(XX, YY)

# create meshgrid for plotting decision boundary
x0, x1 = XX[colx], XX[coly]
linspace_x = np.linspace(x0.min(), x0.max(), 501)
linspace_y = np.linspace(x1.min(), x1.max(), 501)
xx, yy = np.meshgrid(linspace_x, linspace_y)

# make predictions on meshgrid and reshape for plotting
XY = np.c_[xx.ravel(), yy.ravel()]
p = model.predict(XY).reshape(xx.shape)

# plot decision boundary and data points
fig, ax = plt.subplots(1)
ax.contourf(xx, yy, p, cmap="Set1", alpha=0.2)
scatter = ax.scatter(x0, x1, c=YY, cmap="Set1")
ax.set(xlabel=colx, ylabel=coly)

# set title and legend label for current plot
ax.set(title=f"{qualitative_feature} = {key[i]} ")
L = ax.legend(*scatter.legend_elements())

# update legend labels to use top 50 dictionary
for i in range(len(L.get_texts())):
    text = str(L.get_texts()[i])
    new_label = top_50_key[int(text[27])]
    L.get_texts()[i].set_text(new_label)

```

6.0.2 5.1 : Logistic Regression model

6.0.3 5.1.1 : Feature selection

```

[40]: # define a function to compute the score for Logistic Regression model
def cv_score_LRM(cols):
    """
    The funtion compute the score for Logistic Regression model.

    Args:
        cols: a list of column names to use as features in the logistic_
        ↪ regression model.

    Returns:

```

```
A float representing the accuracy score of the logistic regression  
↪model.
```

```
"""
```

```
LR = LogisticRegression(max_iter=150, C = 1.0)  
return cross_val_score(LR, X_train[cols], y_train, cv = 5).mean()
```

```
[41]: # observe Logistic Regression model socre  
LRM_cv_score = get_cv_score(cv_score_LRM)
```

```
Training with columns ['Genre', 'Number of Times Charted', 'Streams']  
The cv score is 0.678
```

```
<class 'numpy.float64'>  
Training with columns ['Chord', 'Number of Times Charted', 'Streams']  
The cv score is 0.678
```

```
<class 'numpy.float64'>  
Training with columns ['Genre', 'Number of Times Charted', 'Artist Followers']  
The cv score is 0.678
```

```
<class 'numpy.float64'>  
Training with columns ['Chord', 'Number of Times Charted', 'Artist Followers']  
The cv score is 0.678
```

```
<class 'numpy.float64'>  
Training with columns ['Genre', 'Number of Times Charted', 'Popularity']  
The cv score is 0.758
```

```
<class 'numpy.float64'>  
Training with columns ['Chord', 'Number of Times Charted', 'Popularity']  
The cv score is 0.762
```

```
<class 'numpy.float64'>  
Training with columns ['Genre', 'Number of Times Charted', 'Danceability']  
The cv score is 0.749
```

```
<class 'numpy.float64'>  
Training with columns ['Chord', 'Number of Times Charted', 'Danceability']  
The cv score is 0.743
```

```
<class 'numpy.float64'>  
Training with columns ['Genre', 'Number of Times Charted', 'Energy']  
The cv score is 0.746
```

```
<class 'numpy.float64'>  
Training with columns ['Chord', 'Number of Times Charted', 'Energy']  
The cv score is 0.748
```

```

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Loudness']
The cv score is 0.747

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Loudness']
The cv score is 0.748

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Speechiness']
The cv score is 0.75

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Speechiness']
The cv score is 0.747

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Acousticness']
The cv score is 0.745

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Acousticness']
The cv score is 0.748

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Liveness']
The cv score is 0.75

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Liveness']
The cv score is 0.748

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Tempo']
The cv score is 0.75

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Tempo']
The cv score is 0.748

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Duration (ms)']
The cv score is 0.709

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Duration (ms)']
The cv score is 0.71

```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Valence']
The cv score is 0.742
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Valence']
The cv score is 0.741
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Artist Followers']
The cv score is 0.678
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Artist Followers']
The cv score is 0.678
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Popularity']
The cv score is 0.678
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Popularity']
The cv score is 0.678
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Danceability']
The cv score is 0.678
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Danceability']
The cv score is 0.678
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Energy']
The cv score is 0.678
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Energy']
The cv score is 0.678
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Loudness']
The cv score is 0.678
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Loudness']
The cv score is 0.678
```

```

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Speechiness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Speechiness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Acousticness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Acousticness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Liveness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Liveness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Duration (ms)']
The cv score is 0.706

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Duration (ms)']
The cv score is 0.706

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Valence']
The cv score is 0.678

```

```

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Popularity']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Popularity']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Danceability']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Danceability']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Energy']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Energy']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Loudness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Loudness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Speechiness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Speechiness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Acousticness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Acousticness']
The cv score is 0.678

```



```

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Liveness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Liveness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Duration (ms)']
The cv score is 0.671

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Duration (ms)']
The cv score is 0.671

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Danceability']
The cv score is 0.707

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Danceability']
The cv score is 0.707

<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Energy']
The cv score is 0.712

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Energy']
The cv score is 0.71

```

```

<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Loudness']
The cv score is 0.714

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Loudness']
The cv score is 0.708

<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Speechiness']
The cv score is 0.717

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Speechiness']
The cv score is 0.708

<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Acousticness']
The cv score is 0.705

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Acousticness']
The cv score is 0.706

<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Liveness']
The cv score is 0.715

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Liveness']
The cv score is 0.71

<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Tempo']
The cv score is 0.713

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Tempo']
The cv score is 0.71

<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Duration (ms)']
The cv score is 0.674

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Duration (ms)']
The cv score is 0.674

```

```

<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Valence']
The cv score is 0.721

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Valence']
The cv score is 0.718

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Energy']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Energy']
The cv score is 0.677

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Loudness']
The cv score is 0.675

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Loudness']
The cv score is 0.675

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Speechiness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Speechiness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Acousticness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Acousticness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Liveness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Liveness']
The cv score is 0.678

```

```

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Loudness']
The cv score is 0.677

<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Loudness']
The cv score is 0.677

<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Speechiness']
The cv score is 0.677

<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Speechiness']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Acousticness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Acousticness']
The cv score is 0.678

```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Liveness']
The cv score is 0.677

<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Liveness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Tempo']
The cv score is 0.677

<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Valence']
The cv score is 0.677

<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Speechiness']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Speechiness']
The cv score is 0.675

<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Acousticness']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Acousticness']
The cv score is 0.676
```

```

<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Liveness']
The cv score is 0.677

<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Liveness']
The cv score is 0.675

<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Tempo']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Tempo']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Valence']
The cv score is 0.677

<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Valence']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Genre', 'Speechiness', 'Acousticness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Speechiness', 'Acousticness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Speechiness', 'Liveness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Speechiness', 'Liveness']
The cv score is 0.678

```

```

<class 'numpy.float64'>
Training with columns ['Genre', 'Speechiness', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Speechiness', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Speechiness', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Speechiness', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Speechiness', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Speechiness', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Acousticness', 'Liveness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Acousticness', 'Liveness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Acousticness', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Acousticness', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Acousticness', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Acousticness', 'Duration (ms)']
The cv score is 0.678

```

```

<class 'numpy.float64'>
Training with columns ['Genre', 'Acousticness', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Acousticness', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Liveness', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Liveness', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Liveness', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Liveness', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Liveness', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Liveness', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Tempo', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Tempo', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Tempo', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Tempo', 'Valence']
The cv score is 0.678

```



```
<class 'numpy.float64'>
Training with columns ['Genre', 'Duration (ms)', 'Valence']
The cv score is 0.678
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Duration (ms)', 'Valence']
The cv score is 0.678
```

```
<class 'numpy.float64'>
```

```
[42]: # observe the best score
display_best_cols(LRM_cv_score)
```

The combination with the highest score is ['Chord', 'Number of Times Charted', 'Popularity'] with a cv score of 0.762.

Based on the analysis and comparison of different feature combinations, the three features ['Chord', 'Number of Times Charted', 'Popularity'] were selected as the best combination for the Logistic Regression model. This combination achieved the highest score, indicating a better predictive performance of the model of around 0.762. In this method we ran our combinations using different pairs of quantitative features and a single qualitative feature 'Chord' at the minimum max iterations that would allow them all to very converge to see what combination would give us the best performance. The combination of ['Chord', 'Number of Times Charted', 'Popularity'] had the highest regression score and so based on these results in addition to the exploratory data analysis, they were chosen as the three features for Logistic Regression model.

```
[43]: # Define the hyperparameter grid
param_grid = {
    'C': [0.001, 0.01, 0.1, 1, 10, 100],
    'penalty': ['l1', 'l2', 'elasticnet', None]
}

# Instantiate a logistic regression model with the desired settings
LR = LogisticRegression(max_iter=1000, solver='saga')

# Instantiate the GridSearchCV object with the logistic regression model and
↳ hyperparameter grid
grid = GridSearchCV(LR, param_grid)

# Fit the GridSearchCV object to the training data, optimizing for the
↳ specified metrics
grid.fit(X_train[['Chord', 'Number of Times Charted', 'Popularity']], y_train)

# Print the best hyperparameters found by GridSearchCV
print(grid.best_params_)
```

```
{'C': 0.1, 'penalty': 'l1'}
```

To have the best number of parameters that can give us the best cv score with the best combinations. We are using GridSearchCV to find it out. The best hyperparameters found by GridSearchCV are printed out using 'grid.best_params_'. The input features for the model are 'Chord', 'Number of Times Charted', and 'Popularity', and the corresponding target variable is 'Y_train'. The best parameters for logistic regression model :{'C': 0.1, 'penalty': 'l1'}

Scaling the data before putting it into a k-NN (k-Nearest Neighbors) machine learning model is important because k-NN is a distance-based algorithm, meaning that it measures the distance between data points to make predictions. If the features in the data are not scaled properly, then the distance between data points may not accurately reflect their true differences. For example, if one feature has a large scale compared to another feature, then the distance between data points will be dominated by the variation in the large-scale feature, which may not be reflective of the actual differences between data points. Scaling the data ensures that all features are on the same scale, so that the distances between data points reflect their true differences in all dimensions. This can improve the accuracy of the k-NN model. Additionally, scaling the data can help to speed up the computation of the k-NN algorithm, since it reduces the range of possible distance values and can make the algorithm more efficient. In our case, our Number of Times Charted could be significantly higher than the Popularity score, which only ranges from 0 to 1.0. Therefore, scaling is necessary

```
[44]: best_combinations_LR = ['Chord', 'Number of Times Charted', 'Popularity']
      # create logistic regression model with best parameters for our data sets

      LR = LogisticRegression(max_iter=1000, C = grid.best_params_['C'], solver = 'saga',
                               penalty = grid.best_params_['penalty'])
      best_score_LR = cross_val_score(LR, X_train[best_combinations_LR], y_train, cv=5).mean()

      print(f"The cv score of parameters {grid.best_params_} and combinations {best_combinations_LR} is {best_score_LR}.")
```

The cv score of parameters {'C': 0.1, 'penalty': 'l1'} and combinations ['Chord', 'Number of Times Charted', 'Popularity'] is 0.7466666666666667.

6.0.4 Apply combination to the model with test set

```
[45]: # # create logistic regression model with best parameters for our data sets
      # LR = LogisticRegression(max_iter=1000, C = 10, solver = 'saga', penalty = 'l1')

      # fit the model
      LR.fit(X_train[best_combinations_LR], y_train)

      lr_train_score = LR.score(X_train[best_combinations_LR], y_train)

      #display the model score
      lr_test_score = LR.score(X_test[best_combinations_LR], y_test)

      print(f"LR CV Score: = {best_score_LR}")
```

```
print(f"LR Train Score: = {lr_train_score}")
print(f"LR Test Score: = {lr_test_score}")
```

LR CV Score: = 0.7466666666666667


LR Train Score: = 0.7475

LR Test Score: = 0.75

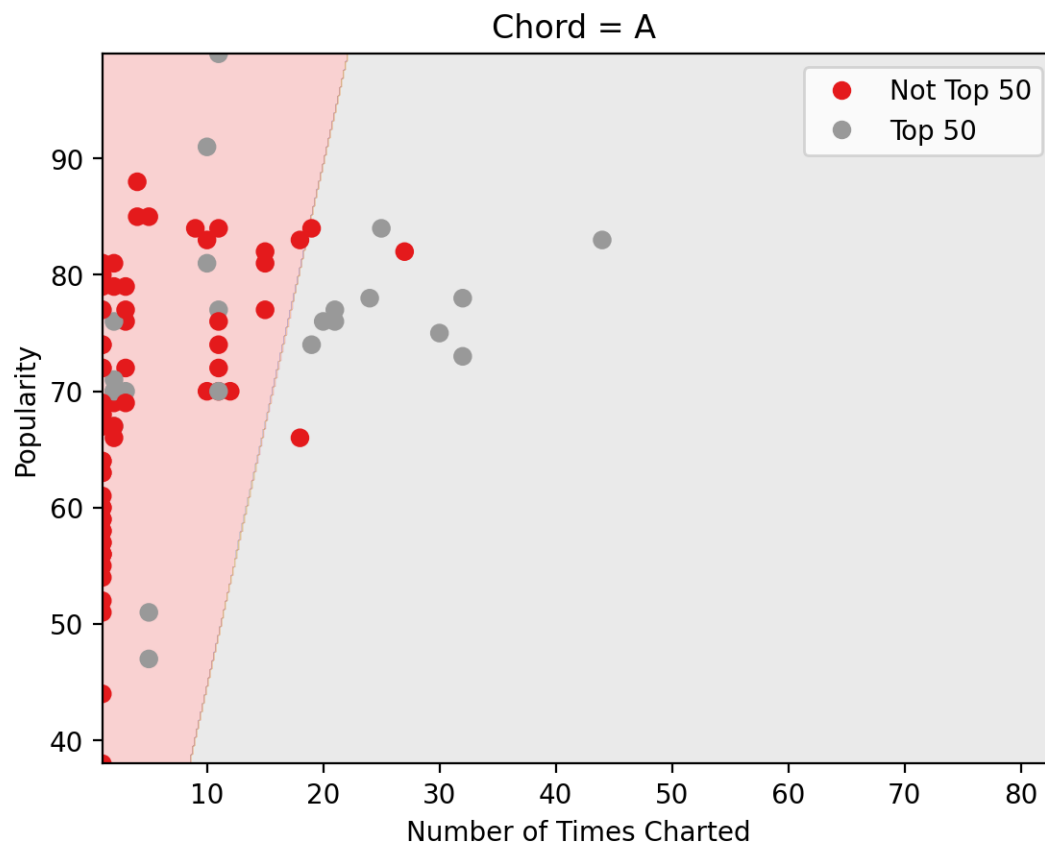
```
[46]: from sklearn.metrics import classification_report
y_pred = LR.predict (X_test[['Chord', 'Number of Times Charted', 'Popularity']])


report = classification_report(y_test, y_pred)
print(report)
confusion_mat = pd.crosstab(y_test, y_pred, rownames=['Actual'], colnames=
    ['Predicted'])
plt.figure (figsize=(3, 3), dpi=100)
sns.heatmap(confusion_mat, annot=True, cmap='Blues', fmt='g')
plt.show()
```


	precision	recall	f1-score	support
False	0.76	0.95	0.84	209
True	0.71	0.30	0.42	91
accuracy			0.75	300
macro avg	0.73	0.62	0.63	300
weighted avg	0.74	0.75	0.71	300




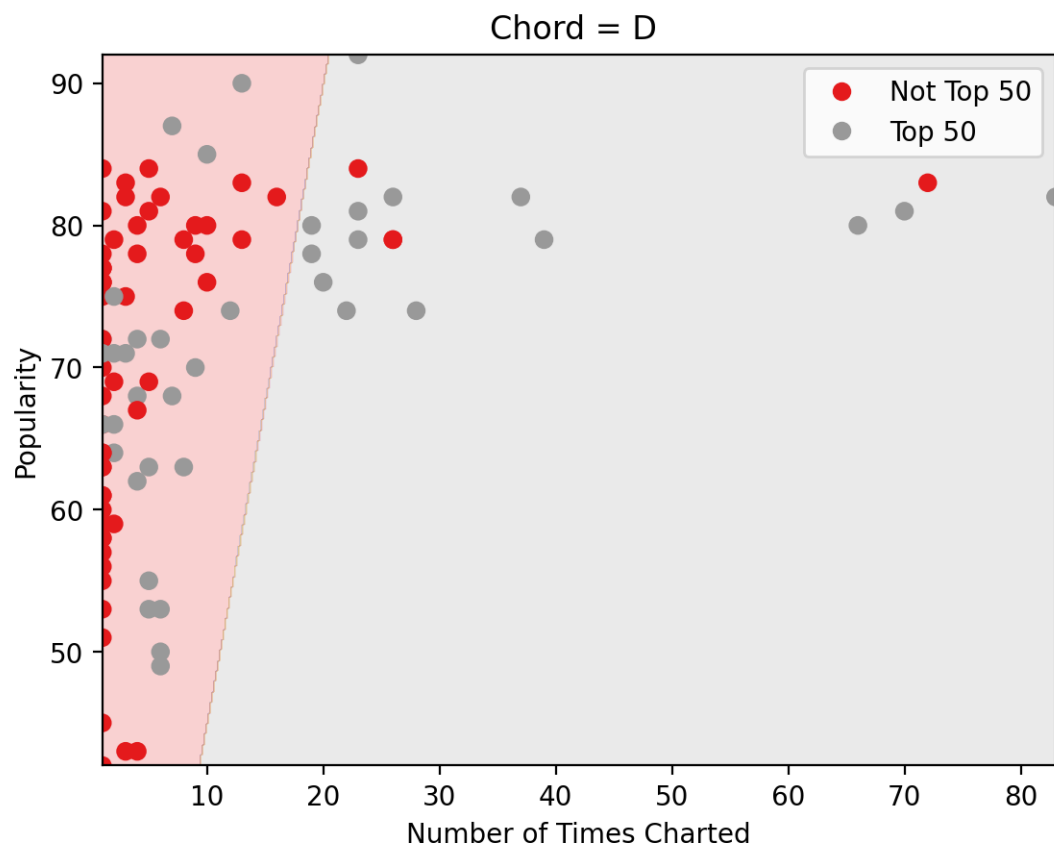
```
plot_decision_regions(LR, X_train, y_train, 'Number of Times Charted',  
                    ↪ 'Popularity', 'Chord')
```

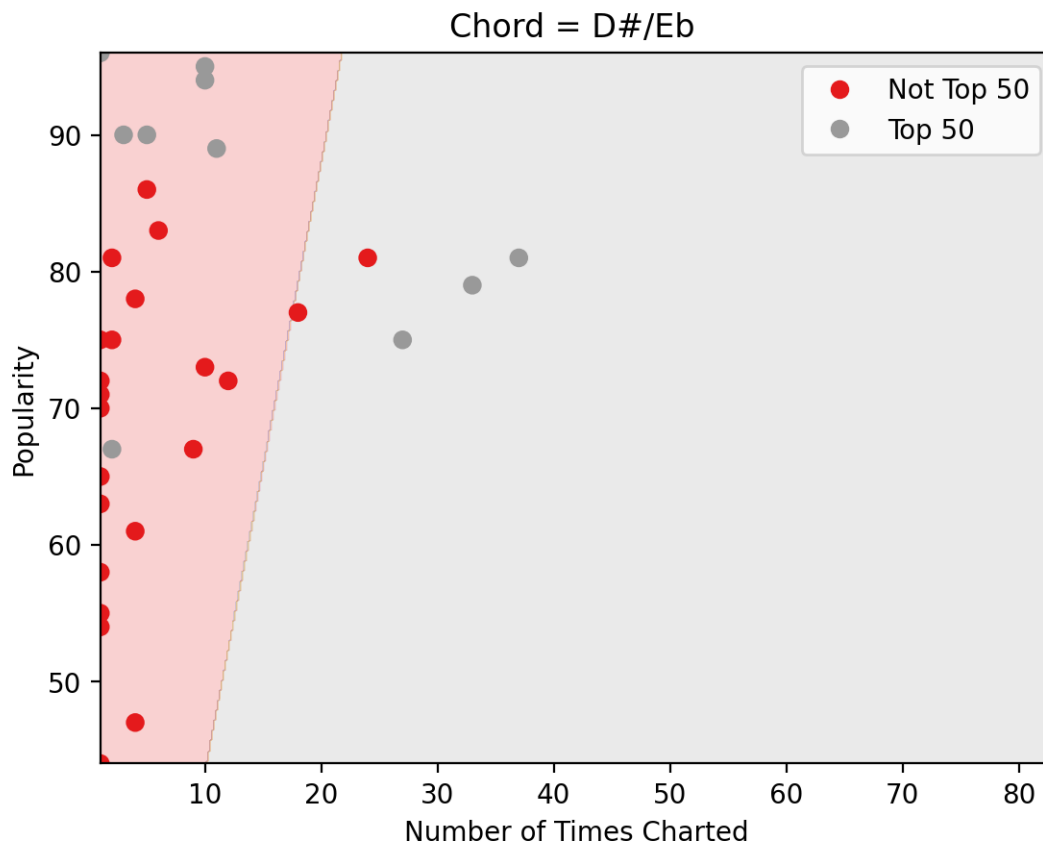
























The logistic regression model achieved an accuracy score of 0.75 on the test data, which means it was able to correctly predict the class label for 75% of the samples in the test set. However, the confusion matrix shows that the model struggled to predict the True class label, as evidenced by a recall score of only 0.27. This indicates that the model is biased towards predicting the False class label and has difficulty identifying samples that belong to the True class.

When we examine the decision regions plotted, it becomes apparent that the model might be overgeneralizing the negative class, leading to the lower recall score for the positive class. This issue is particularly noticeable in the graphs for chords G#, F, and C#. In these graphs, we see that the logistic model calculated the positive class region to be quite small and predominantly restricted to the upper end of the feature space. However, a higher concentration of positive examples actually exists at the lower end.

This example highlights some of the errors in the logistic regression model. Based on our prior data analysis, we know that the ‘Number of Times Charted’ variable was relatively positively correlated with whether a song was in the top 50. However, in these decision regions, the boundary line not only lacks logical sense (such as with G# where the area for Top 50 is nearly invisible in the right bottom corner) but also fails to accurately predict the boolean.

In summary, the logistic regression model is not well-suited for capturing the complex relationships between the features and the target variable in this dataset. As a result, it should not be used to determine if a sample belongs to the Top 50 category, as it is biased towards predicting negative

outcomes.

6.0.5 Random Forests Model

```
[48]: # define a function to compute the score for Random Forests Model
def score_RFM(cols):
    """
    The funtion compute the score for Random Forests Model.

    Args:
        cols: a list of column names to use as features in the Random Forests_
        ↪Model.

    Returns:
        A float representing the accuracy score of the Random Forests Model.
    """
    RF = RandomForestClassifier(n_estimators=100)
    return cross_val_score(RF, X_train[cols], y_train, cv = 5).mean()

[49]: # observe Random Forests Model socre
RFM_score = get_cv_score(score_RFM)
```

```
Training with columns ['Genre', 'Number of Times Charted', 'Streams']
The cv score is 0.775
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Streams']
The cv score is 0.774
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Artist Followers']
The cv score is 0.795
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Artist Followers']
The cv score is 0.757
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Popularity']
The cv score is 0.753
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Popularity']
The cv score is 0.735
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Danceability']
The cv score is 0.737
```



```

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Danceability']
The cv score is 0.694

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Energy']
The cv score is 0.728

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Energy']
The cv score is 0.718

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Loudness']
The cv score is 0.726

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Loudness']
The cv score is 0.709

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Speechiness']
The cv score is 0.722

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Speechiness']
The cv score is 0.712

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Acousticness']
The cv score is 0.728

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Acousticness']
The cv score is 0.696

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Liveness']
The cv score is 0.718

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Liveness']
The cv score is 0.7

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Tempo']
The cv score is 0.713

```

```

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Tempo']
The cv score is 0.709

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Duration (ms)']
The cv score is 0.703

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Duration (ms)']
The cv score is 0.704

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Valence']
The cv score is 0.712

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Valence']
The cv score is 0.718

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Artist Followers']
The cv score is 0.712

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Artist Followers']
The cv score is 0.701

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Popularity']
The cv score is 0.709

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Popularity']
The cv score is 0.703

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Danceability']
The cv score is 0.662

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Danceability']
The cv score is 0.673

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Energy']
The cv score is 0.691

```

```

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Energy']
The cv score is 0.677

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Loudness']
The cv score is 0.667

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Loudness']
The cv score is 0.658

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Speechiness']
The cv score is 0.669

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Speechiness']
The cv score is 0.662

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Acousticness']
The cv score is 0.664

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Acousticness']
The cv score is 0.648

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Liveness']
The cv score is 0.662

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Liveness']
The cv score is 0.679

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Tempo']
The cv score is 0.68

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Tempo']
The cv score is 0.662

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Duration (ms)']
The cv score is 0.665

```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Duration (ms)']
The cv score is 0.656

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Valence']
The cv score is 0.679

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Valence']
The cv score is 0.658

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Popularity']
The cv score is 0.73

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Popularity']
The cv score is 0.709

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Danceability']
The cv score is 0.644

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Danceability']
The cv score is 0.633

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Energy']
The cv score is 0.664

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Energy']
The cv score is 0.652

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Loudness']
The cv score is 0.702

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Loudness']
The cv score is 0.663

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Speechiness']
The cv score is 0.652
```

```

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Speechiness']
The cv score is 0.639

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Acousticness']
The cv score is 0.657

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Acousticness']
The cv score is 0.643

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Liveness']
The cv score is 0.652

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Liveness']
The cv score is 0.645

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Tempo']
The cv score is 0.653

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Tempo']
The cv score is 0.663

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Duration (ms)']
The cv score is 0.662

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Duration (ms)']
The cv score is 0.646

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Valence']
The cv score is 0.662

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Valence']
The cv score is 0.638

<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Danceability']
The cv score is 0.659

```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Danceability']
The cv score is 0.671
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Energy']
The cv score is 0.673
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Energy']
The cv score is 0.67
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Loudness']
The cv score is 0.658
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Loudness']
The cv score is 0.649
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Speechiness']
The cv score is 0.648
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Speechiness']
The cv score is 0.667
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Acousticness']
The cv score is 0.651
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Acousticness']
The cv score is 0.672
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Liveness']
The cv score is 0.69
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Liveness']
The cv score is 0.667
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Tempo']
The cv score is 0.68
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Tempo']
The cv score is 0.682

<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Duration (ms)']
The cv score is 0.668

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Duration (ms)']
The cv score is 0.652

<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Valence']
The cv score is 0.67

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Valence']
The cv score is 0.663

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Energy']
The cv score is 0.641

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Energy']
The cv score is 0.595

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Loudness']
The cv score is 0.647

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Loudness']
The cv score is 0.628

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Speechiness']
The cv score is 0.628

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Speechiness']
The cv score is 0.632

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Acousticness']
The cv score is 0.614
```

```

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Acousticness']
The cv score is 0.629

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Liveness']
The cv score is 0.63

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Liveness']
The cv score is 0.611

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Tempo']
The cv score is 0.622

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Tempo']
The cv score is 0.621

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Duration (ms)']
The cv score is 0.604

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Duration (ms)']
The cv score is 0.621

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Valence']
The cv score is 0.63

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Valence']
The cv score is 0.602

<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Loudness']
The cv score is 0.641

<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Loudness']
The cv score is 0.639

<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Speechiness']
The cv score is 0.639

```



```
<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Speechiness']
The cv score is 0.628

<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Acousticness']
The cv score is 0.611

<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Acousticness']
The cv score is 0.637

<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Liveness']
The cv score is 0.632

<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Liveness']
The cv score is 0.628

<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Tempo']
The cv score is 0.638

<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Tempo']
The cv score is 0.638

<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Duration (ms)']
The cv score is 0.639

<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Duration (ms)']
The cv score is 0.617

<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Valence']
The cv score is 0.639

<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Valence']
The cv score is 0.638

<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Speechiness']
The cv score is 0.636
```

```

<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Speechiness']
The cv score is 0.633

<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Acousticness']
The cv score is 0.635

<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Acousticness']
The cv score is 0.622

<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Liveness']
The cv score is 0.637

<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Liveness']
The cv score is 0.638

<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Tempo']
The cv score is 0.628

<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Tempo']
The cv score is 0.634

<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Duration (ms)']
The cv score is 0.656

<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Duration (ms)']
The cv score is 0.602

<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Valence']
The cv score is 0.634

<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Valence']
The cv score is 0.606

<class 'numpy.float64'>
Training with columns ['Genre', 'Speechiness', 'Acousticness']
The cv score is 0.635

```

```

<class 'numpy.float64'>
Training with columns ['Chord', 'Speechiness', 'Acousticness']
The cv score is 0.613

<class 'numpy.float64'>
Training with columns ['Genre', 'Speechiness', 'Liveness']
The cv score is 0.606

<class 'numpy.float64'>
Training with columns ['Chord', 'Speechiness', 'Liveness']
The cv score is 0.614

<class 'numpy.float64'>
Training with columns ['Genre', 'Speechiness', 'Tempo']
The cv score is 0.631

<class 'numpy.float64'>
Training with columns ['Chord', 'Speechiness', 'Tempo']
The cv score is 0.628

<class 'numpy.float64'>
Training with columns ['Genre', 'Speechiness', 'Duration (ms)']
The cv score is 0.614

<class 'numpy.float64'>
Training with columns ['Chord', 'Speechiness', 'Duration (ms)']
The cv score is 0.611

<class 'numpy.float64'>
Training with columns ['Genre', 'Speechiness', 'Valence']
The cv score is 0.619

<class 'numpy.float64'>
Training with columns ['Chord', 'Speechiness', 'Valence']
The cv score is 0.601

<class 'numpy.float64'>
Training with columns ['Genre', 'Acousticness', 'Liveness']
The cv score is 0.615

<class 'numpy.float64'>
Training with columns ['Chord', 'Acousticness', 'Liveness']
The cv score is 0.628

<class 'numpy.float64'>
Training with columns ['Genre', 'Acousticness', 'Tempo']
The cv score is 0.625

```

```

<class 'numpy.float64'>
Training with columns ['Chord', 'Acousticness', 'Tempo']
The cv score is 0.609

<class 'numpy.float64'>
Training with columns ['Genre', 'Acousticness', 'Duration (ms)']
The cv score is 0.628

<class 'numpy.float64'>
Training with columns ['Chord', 'Acousticness', 'Duration (ms)']
The cv score is 0.618

<class 'numpy.float64'>
Training with columns ['Genre', 'Acousticness', 'Valence']
The cv score is 0.617

<class 'numpy.float64'>
Training with columns ['Chord', 'Acousticness', 'Valence']
The cv score is 0.618

<class 'numpy.float64'>
Training with columns ['Genre', 'Liveness', 'Tempo']
The cv score is 0.646

<class 'numpy.float64'>
Training with columns ['Chord', 'Liveness', 'Tempo']
The cv score is 0.62

<class 'numpy.float64'>
Training with columns ['Genre', 'Liveness', 'Duration (ms)']
The cv score is 0.625

<class 'numpy.float64'>
Training with columns ['Chord', 'Liveness', 'Duration (ms)']
The cv score is 0.616

<class 'numpy.float64'>
Training with columns ['Genre', 'Liveness', 'Valence']
The cv score is 0.631

<class 'numpy.float64'>
Training with columns ['Chord', 'Liveness', 'Valence']
The cv score is 0.622

<class 'numpy.float64'>
Training with columns ['Genre', 'Tempo', 'Duration (ms)']
The cv score is 0.656

```

```

<class 'numpy.float64'>
Training with columns ['Chord', 'Tempo', 'Duration (ms)']
The cv score is 0.633

<class 'numpy.float64'>
Training with columns ['Genre', 'Tempo', 'Valence']
The cv score is 0.636

<class 'numpy.float64'>
Training with columns ['Chord', 'Tempo', 'Valence']
The cv score is 0.633

<class 'numpy.float64'>
Training with columns ['Genre', 'Duration (ms)', 'Valence']
The cv score is 0.629

<class 'numpy.float64'>
Training with columns ['Chord', 'Duration (ms)', 'Valence']
The cv score is 0.604

<class 'numpy.float64'>

```

```

[50]: # observe the best score for Forests Model score
      display_best_cols(RFM_score)

```

The combination with the highest score is ['Genre', 'Number of Times Charted', 'Artist Followers'] with a cv score of 0.795.

Based on the results of the feature selection process using the Random Forests Model, the combination of ['Genre', 'Number of Times Charted', 'Artist Followers'] was found to have the highest score of around 0.79.

The selection of these features was based on the exploratory data analysis performed on the dataset, which showed that these features have a relatively strong correlation with the target variable (Hit or Flop).

Therefore, based on the results of the feature selection process and the exploratory data analysis, the combination of ['Genre', 'Number of Times Charted', 'Artist Followers'] was chosen as the most promising set of features to use for the Random Forests Model.

6.0.6 Cross-validation to choose complexity parameters (max_depth, n_estimators)

Part 1: max_depth Selection

```

[51]: # Create figure and axes
      fig, ax = plt.subplots(1)

      # Create "best" variables
      best_max_depth = 0

```

```

best_max_depth_cv = 0
best_combinations_RF = ['Genre', 'Number of Times Charted', 'Streams']

# Run cv function for max_depth selection from 1 to 30
for i in range(1,30):
    RF = RandomForestClassifier(n_estimators = 100 , max_depth = i)
    cv = cross_val_score(RF, X_train[best_combinations_RF], y_train, cv = 5).
    ↪mean()

    # Create scatter plot to visualize results
    ax.scatter(i, cv, color = 'blue')

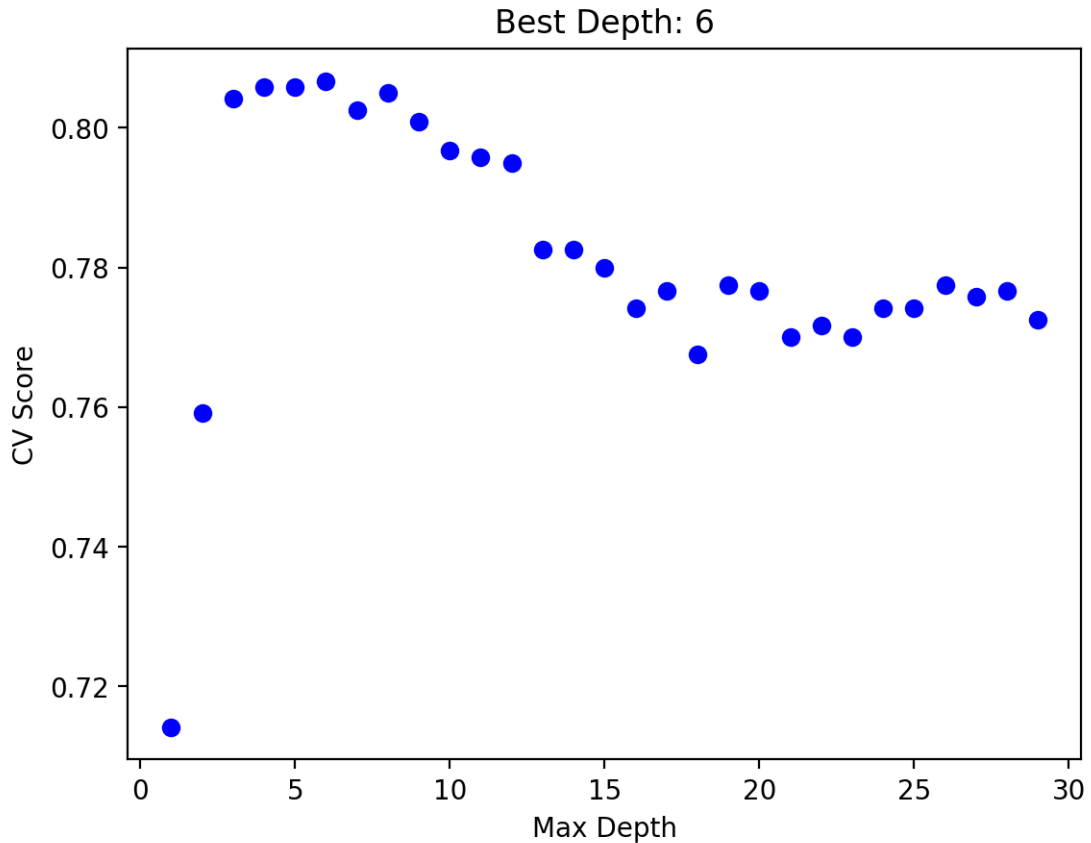
    # Determine best max_depth and its corresponding cv
    if cv > best_max_depth_cv:
        best_max_depth_cv = cv
        best_max_depth = i

# Printing best max_depth and its corresponding cv
print(f"Best max depth is {best_max_depth} with a cv of {best_max_depth_cv}")

# Label scatterplot
l = ax.set(title = "Best Depth: " + str(best_max_depth), xlabel = "Max Depth", ↪
    ↪ylabel = "CV Score")

```

Best max depth is 6 with a cv of 0.8066666666666666



Part 2: n_estimator Selection

```
[52]: # Create figure and axes
fig, ax = plt.subplots(1)

# Create "best" variables
best_n_estimators = 0
best_n_estimators_cv = 0
best_combinations_RF = ['Genre', 'Number of Times Charted', 'Streams']

for i in range(1,400,10):
    RF = RandomForestClassifier(n_estimators = i)
    cv = cross_val_score(RF, X_train[best_combinations_RF], y_train, cv = 5).
    ↪mean()

    # Create scatter plot to visualize results
    ax.scatter(i, cv, color = 'blue')

    # Determine best max_depth and its corresponding cv
    if cv > best_n_estimators_cv:
```

```

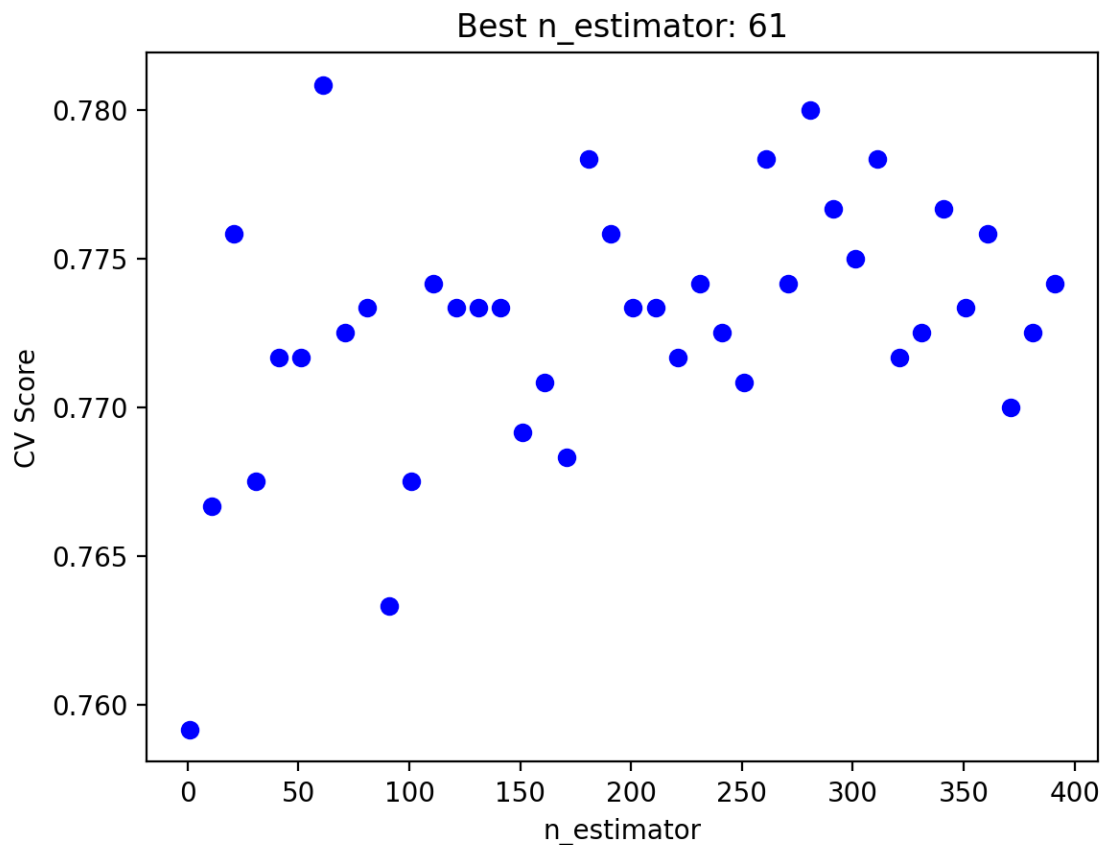
best_n_estimators_cv = cv
best_n_estimators = i

# Printing best max_depth and its corresponding cv
print(f"Best n_estimator is {best_n_estimators} with a cv of_
↪{best_n_estimators_cv}")

# Label scatterplot
l = ax.set(title = "Best n_estimator: " + str(best_n_estimators), xlabel =_
↪"n_estimator", ylabel = "CV Score")

```

Best n_estimator is 61 with a cv of 0.7808333333333334



However, we observed that if we recalculated `n_estimator`, based on the `max_depth` score that we calculated in the previous step, we can get a `n_estimator` score that ultimately results in a higher cv that we can use for our random forest, which is why we decided that we should not tune `max_depth` and `n_estimators` separately because they result in a lower score (i.e. we evaluate `max_depth` and then hold it constant to evaluate `n_estimators` as shown instead below)


```

[53]: # Create figure and axes
fig, ax = plt.subplots(1)

# Create "best" variables
best_n_estimators = 0
best_n_estimators_cv = 0
best_combinations_RF = ['Genre', 'Number of Times Charted', 'Streams']

for i in range(1,400,10):
    RF = RandomForestClassifier(n_estimators = i, max_depth = best_max_depth)
    cv = cross_val_score(RF, X_train[best_combinations_RF], y_train, cv = 5).
    ↪mean()

    # Create scatter plot to visualize results
    ax.scatter(i, cv, color = 'blue')

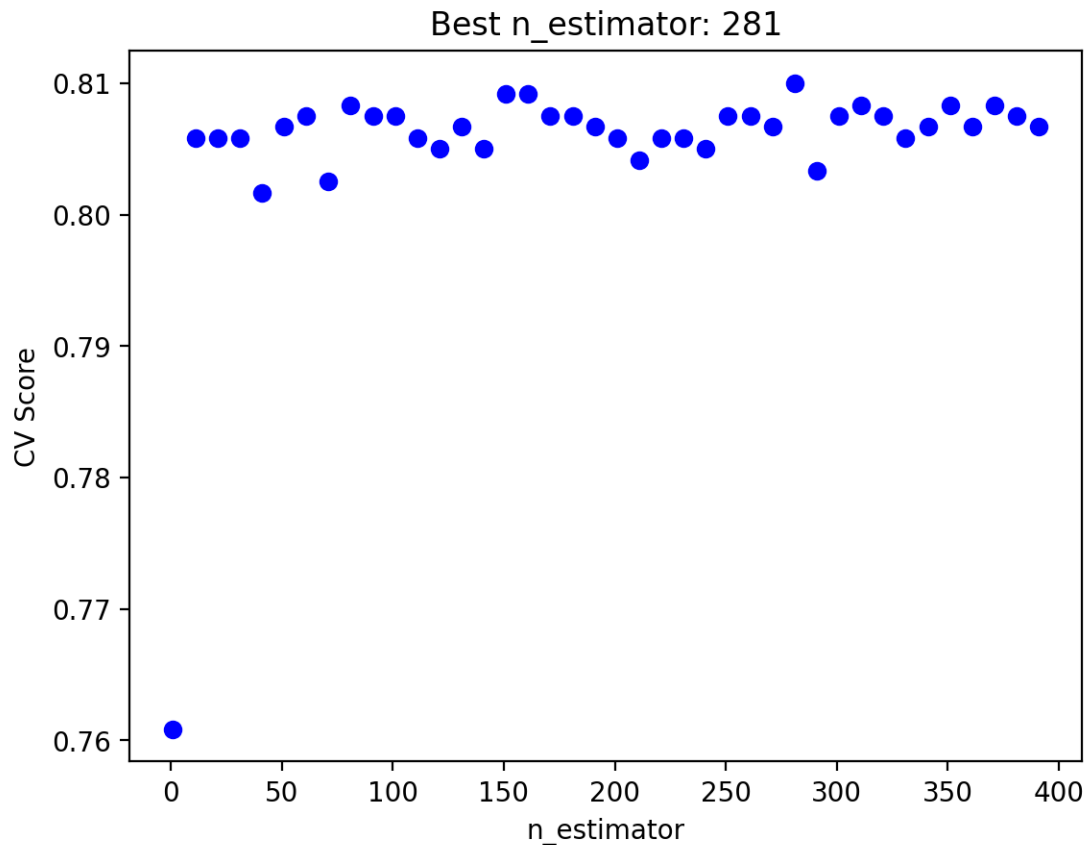
    # Determine best max_depth and its corresponding cv
    if cv > best_n_estimators_cv:
        best_n_estimators_cv = cv
        best_n_estimators = i

# Printing best max_depth and its corresponding cv
print(f"Best n_estimator is {best_n_estimators} with a cv of_
    ↪{best_n_estimators_cv}")

# Label scatterplot
l = ax.set(title = "Best n_estimator: " + str(best_n_estimators), xlabel =_
    ↪"n_estimator", ylabel = "CV Score")

```

Best n_estimator is 281 with a cv of 0.8099999999999999



```
[54]: # create Random Forest model with best parameters for our data sets
RF = RandomForestClassifier(n_estimators = best_n_estimators, max_depth = ␣
↪best_max_depth)

# fit the model
RF.fit(X_train[best_combinations_RF], y_train)

#display the model score
rf_train_score = RF.score(X_train[best_combinations_RF], y_train)

rf_test_score = RF.score(X_test[best_combinations_RF], y_test)

print(f"RF CV Score: = {best_n_estimators_cv}")
print(f"RF Train Score: = {rf_train_score}")
print(f"RF Test Score: = {rf_test_score}")
```

```
RF CV Score: = 0.8099999999999999
RF Train Score: = 0.8325
RF Test Score: = 0.8433333333333334
```

```
[55]: import numpy as np
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import RandomizedSearchCV

# Define the hyperparameter grid
param_dist = {
    'n_estimators': np.arange(1, 1001, 10),
    'max_depth': np.arange(1, 101, 10)
}

# Instantiate a Random Forests Model with the desired settings
RF = RandomForestClassifier()

# Instantiate the RandomizedSearchCV object with the Random Forest model,
↳ hyperparameter grid, and number of iterations
random_search = RandomizedSearchCV(RF, param_distributions=param_dist,
↳ n_iter=50, cv=5, random_state=42)

# Fit the RandomizedSearchCV object to the training data, optimizing for the
↳ specified metrics
random_search.fit(X_train[['Genre', 'Number of Times Charted', 'Streams']],
↳ y_train)

# Print the best hyperparameters found by RandomizedSearchCV
print(random_search.best_params_)
```

```
{'n_estimators': 361, 'max_depth': 11}
```

```
[56]: best_combinations_RF = ['Genre', 'Number of Times Charted', 'Streams']
RF = RandomForestClassifier(n_estimators = random_search.
↳ best_params_['n_estimators'] , max_depth = random_search.
↳ best_params_['max_depth'])
best_score_RF = cross_val_score(RF, X_train[best_combinations_RF], y_train, cv=
↳ 5).mean()

print(f"The cv score of parameters {random_search.best_params_} and
↳ combinations {best_combinations_RF} is {best_score_RF}.")
```

The cv score of parameters {'n_estimators': 361, 'max_depth': 11} and combinations ['Genre', 'Number of Times Charted', 'Streams'] is 0.7949999999999999.

6.0.7 Apply the best combination into test set of Random Forest model

```
[57]: # create Random Forest model with best parameters for our data sets
RF = RandomForestClassifier(n_estimators = random_search.
    ↪best_params_['n_estimators'], max_depth = random_search.
    ↪best_params_['max_depth'])

# fit the model
RF.fit(X_train[best_combinations_RF], y_train)

#display the model score
rf_train_score = RF.score(X_train[best_combinations_RF], y_train)

rf_test_score = RF.score(X_test[best_combinations_RF], y_test)

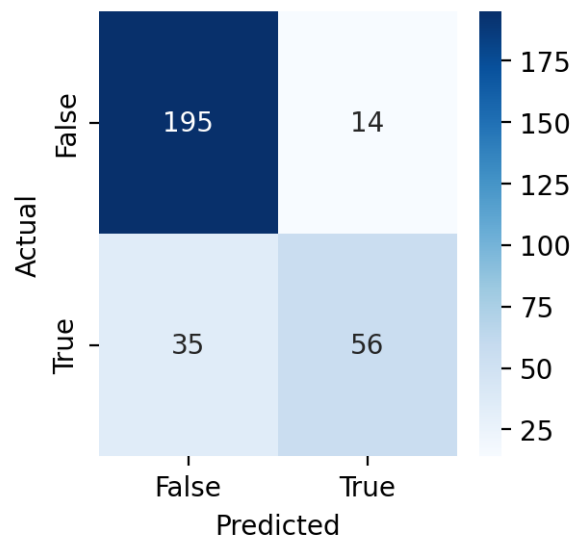
print(f"RF CV Score: = {best_score_RF}")
print(f"RF Train Score: = {rf_train_score}")
print(f"RF Test Score: = {rf_test_score}")
```

```
RF CV Score: = 0.7949999999999999
RF Train Score: = 0.9325
RF Test Score: = 0.8366666666666667
```


```
[58]: from sklearn.metrics import classification_report
y_pred = RF.predict(X_test[best_combinations_RF])


report = classification_report(y_test, y_pred)
print(report)
confusion_mat = pd.crosstab(y_test, y_pred, rownames=['Actual'], colnames=
    ↪['Predicted'])
plt.figure(figsize=(3, 3), dpi=100)
sns.heatmap(confusion_mat, annot=True, cmap='Blues', fmt='g')
plt.show()
```

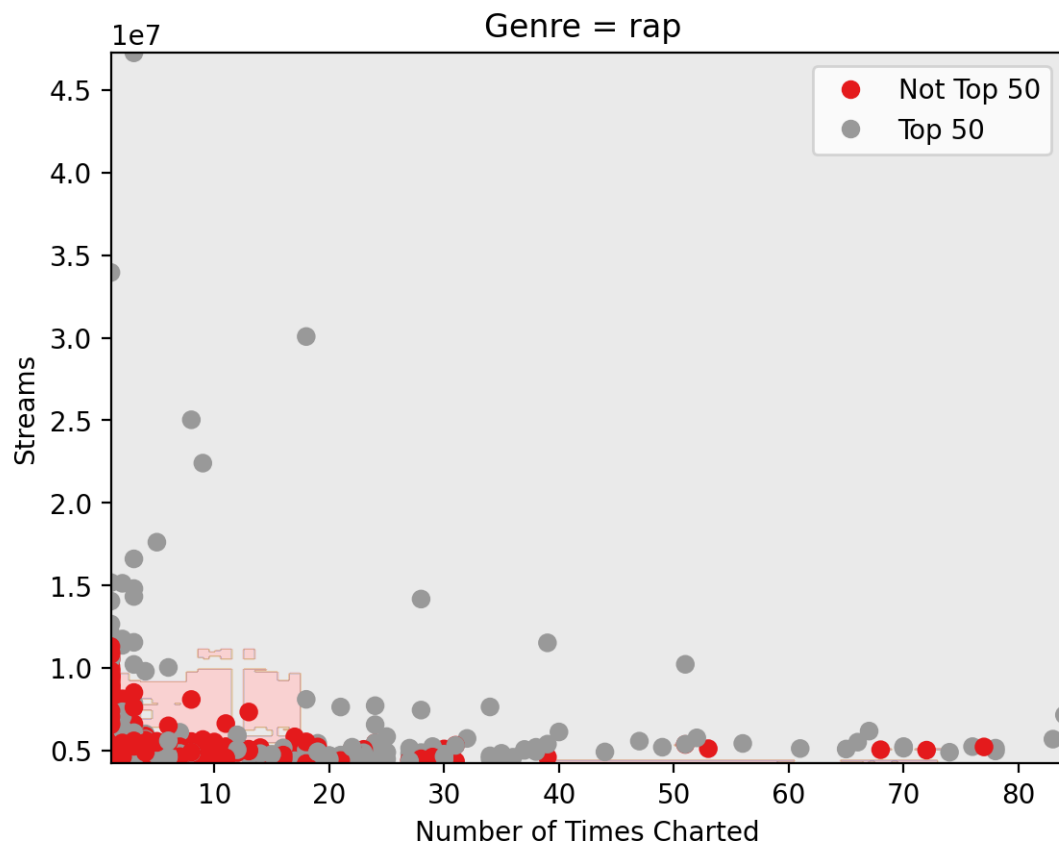
	precision	recall	f1-score	support
False	0.85	0.93	0.89	209
True	0.80	0.62	0.70	91
accuracy			0.84	300
macro avg	0.82	0.77	0.79	300
weighted avg	0.83	0.84	0.83	300

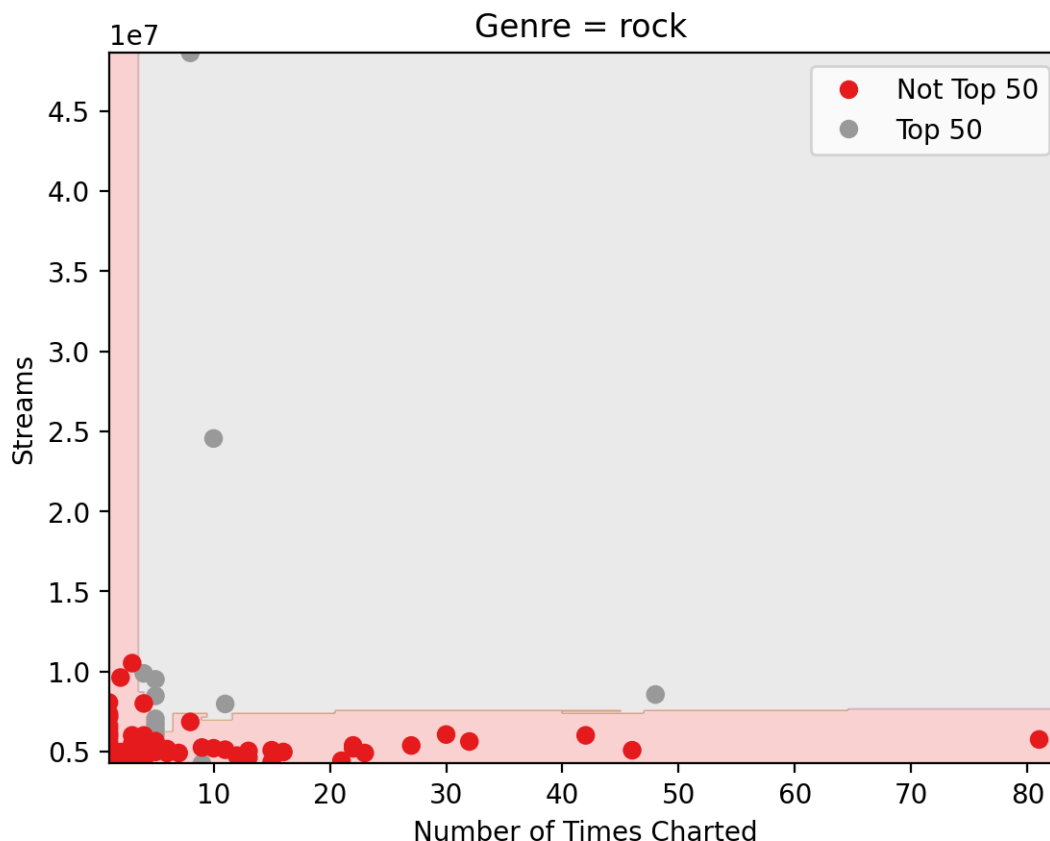


```
[59]: plot_decision_regions(RF, X_train, y_train, 'Number of Times Charted', 'Streams', 'Genre')
```









Analysis

The Random Forest (RF) model exhibits high accuracy on the training data with a score of 0.9308, which means it correctly classified 93.08% of the samples. However, the model's performance on the test data is slightly lower, with an accuracy of 0.8267 (82.67%). The difference between the train and test scores may suggest overfitting, where the model captures noise in the data instead of the true underlying patterns. The decision regions, particularly for rap and pop genres, appear to be very rigid and specific, which further indicates overfitting. In contrast, the decision regions for rock and R&B genres, which have fewer positive data points, are less specific and more flexible. The hyper-specific and rigidly defined regions of the decision plot might lead to poor performance on larger datasets, as the model could struggle to generalize patterns beyond the specific values it was trained on. For example, while the features of streams and times charted may be good indicators of whether a song is in the Top 50, relying on the exact combination of stream and chart count as a reliable indicator of a song's success may be overly simplistic. Additionally, the boundaries in the decision plot may not be applicable outside of this specific dataset, especially in cases where the model has captured noise in the data rather than true patterns. In conclusion, while the RF model performed exceptionally well on the training data, it is important to be cautious when applying it to larger and more diverse datasets. The model's hyper-specific decision regions and potential overfitting could limit its ability to generalize patterns beyond the specific values it was trained on.

6.0.8 K-Nearest Neighbours (KNN) Models

```
[60]: from sklearn.neighbors import KNeighborsClassifier
```

```
knn = KNeighborsClassifier(n_neighbors = 3)
knn.fit(X_train, y_train)
```

```
[60]: KNeighborsClassifier(n_neighbors=3)
```

```
[61]: from sklearn.metrics import classification_report
```

```
y_pred = knn.predict(X_test)
```

```
report = classification_report(y_test, y_pred)
```

```
print(report)
```

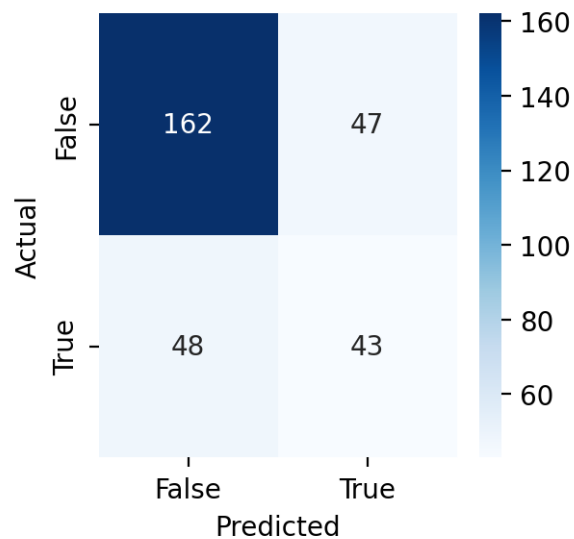
```
confusion_mat = pd.crosstab(y_test, y_pred, rownames=['Actual'],
                             colnames=['Predicted'])
```

```
plt.figure(figsize=(3, 3), dpi=100)
```

```
sns.heatmap(confusion_mat, annot=True, cmap='Blues', fmt='g')
```

```
plt.show()
```

		precision	recall	f1-score	support
	False	0.77	0.78	0.77	209
	True	0.48	0.47	0.48	91
accuracy				0.68	300
macro avg		0.62	0.62	0.62	300
weighted avg		0.68	0.68	0.68	300



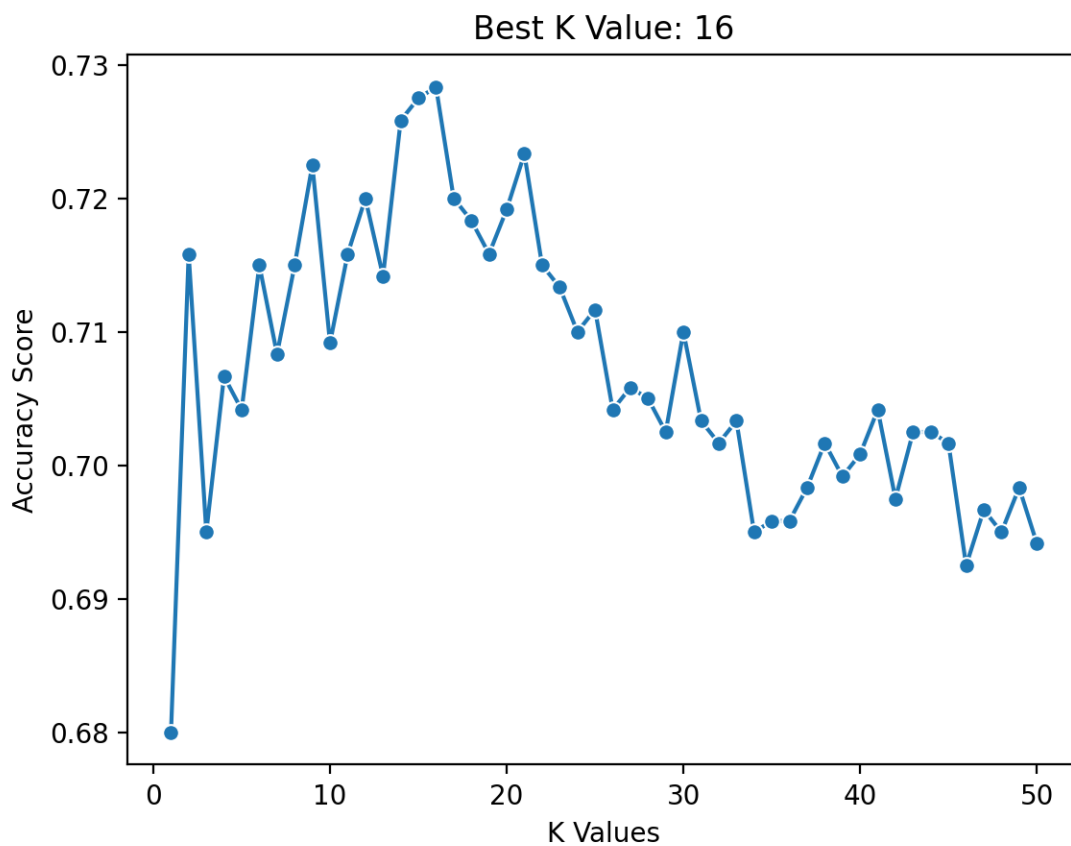
```
[62]: k_values = [i for i in range (1,51)]
      scores = []

      for k in k_values:
          knn = KNeighborsClassifier(n_neighbors=k)
          score = cross_val_score(knn, X_train, y_train, cv=5)
          scores.append(np.mean(score))

      best_k = k_values[np.argmax(scores)]
      best_k

      sns.lineplot(x = k_values, y = scores, marker = 'o')
      plt.xlabel("K Values")
      plt.ylabel("Accuracy Score")
      plt.title(f"Best K Value: {best_k}")
```

```
[62]: Text(0.5, 1.0, 'Best K Value: 16')
```



```
[63]: # define a function to compute the score for KNN
      def score_KNN(cols):
```

```

"""
The function compute the score for Random Forests Model.

Args:
    cols: a list of column names to use as features in the Random Forests
    Model.

Returns:
    A float representing the accuracy score of the Random Forests Model.
"""
knn = KNeighborsClassifier(n_neighbors = best_k)

return cross_val_score(knn, X_train[cols], y_train, cv=5).mean()

```

```
[64]: KNN_score = get_cv_score(score_KNN)
```

```

Training with columns ['Genre', 'Number of Times Charted', 'Streams']
The cv score is 0.709

```

```

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Streams']
The cv score is 0.709

```

```

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Artist Followers']
The cv score is 0.707

```

```

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Artist Followers']
The cv score is 0.7

```

```

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Popularity']
The cv score is 0.784

```

```

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Popularity']
The cv score is 0.783

```

```

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Danceability']
The cv score is 0.753

```

```

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Danceability']
The cv score is 0.757

```

```

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Energy']
The cv score is 0.756

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Energy']
The cv score is 0.759

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Loudness']
The cv score is 0.758

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Loudness']
The cv score is 0.756

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Speechiness']
The cv score is 0.756

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Speechiness']
The cv score is 0.755

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Acousticness']
The cv score is 0.748

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Acousticness']
The cv score is 0.757

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Liveness']
The cv score is 0.757

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Liveness']
The cv score is 0.751

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Tempo']
The cv score is 0.753

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Tempo']
The cv score is 0.749

```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Duration (ms)']
The cv score is 0.649
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Duration (ms)']
The cv score is 0.649
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Valence']
The cv score is 0.748
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Valence']
The cv score is 0.756
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Artist Followers']
The cv score is 0.729
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Artist Followers']
The cv score is 0.729
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Popularity']
The cv score is 0.709
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Popularity']
The cv score is 0.709
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Danceability']
The cv score is 0.709
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Danceability']
The cv score is 0.709
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Energy']
The cv score is 0.709
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Energy']
The cv score is 0.709
```

```

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Loudness']
The cv score is 0.709

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Loudness']
The cv score is 0.709

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Speechiness']
The cv score is 0.709

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Speechiness']
The cv score is 0.709

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Acousticness']
The cv score is 0.709

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Acousticness']
The cv score is 0.709

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Liveness']
The cv score is 0.709

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Liveness']
The cv score is 0.709

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Tempo']
The cv score is 0.709

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Tempo']
The cv score is 0.709

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Duration (ms)']
The cv score is 0.711

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Duration (ms)']
The cv score is 0.711

```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Valence']
The cv score is 0.709

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Valence']
The cv score is 0.709

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Popularity']
The cv score is 0.693

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Popularity']
The cv score is 0.695

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Danceability']
The cv score is 0.682

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Danceability']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Energy']
The cv score is 0.684

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Energy']
The cv score is 0.68

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Loudness']
The cv score is 0.692

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Loudness']
The cv score is 0.689

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Speechiness']
The cv score is 0.684

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Speechiness']
The cv score is 0.677
```

```

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Acousticness']
The cv score is 0.69

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Acousticness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Liveness']
The cv score is 0.677

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Liveness']
The cv score is 0.679

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Tempo']
The cv score is 0.682

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Tempo']
The cv score is 0.682

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Duration (ms)']
The cv score is 0.675

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Duration (ms)']
The cv score is 0.675

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Valence']
The cv score is 0.677

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Valence']
The cv score is 0.677

<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Danceability']
The cv score is 0.702

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Danceability']
The cv score is 0.701

```



```
<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Energy']
The cv score is 0.706

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Energy']
The cv score is 0.694

<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Loudness']
The cv score is 0.706

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Loudness']
The cv score is 0.708

<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Speechiness']
The cv score is 0.703

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Speechiness']
The cv score is 0.706

<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Acousticness']
The cv score is 0.713

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Acousticness']
The cv score is 0.697

<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Liveness']
The cv score is 0.713

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Liveness']
The cv score is 0.705

<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Tempo']
The cv score is 0.698

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Tempo']
The cv score is 0.7
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Duration (ms)']
The cv score is 0.649

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Duration (ms)']
The cv score is 0.649

<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Valence']
The cv score is 0.705

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Valence']
The cv score is 0.698

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Energy']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Energy']
The cv score is 0.662

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Loudness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Loudness']
The cv score is 0.665

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Speechiness']
The cv score is 0.668

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Speechiness']
The cv score is 0.654

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Acousticness']
The cv score is 0.667

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Acousticness']
The cv score is 0.669
```

```

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Liveness']
The cv score is 0.658

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Liveness']
The cv score is 0.663

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Tempo']
The cv score is 0.662

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Tempo']
The cv score is 0.671

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Duration (ms)']
The cv score is 0.648

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Duration (ms)']
The cv score is 0.648

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Valence']
The cv score is 0.667

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Valence']
The cv score is 0.662

<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Loudness']
The cv score is 0.67

<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Loudness']
The cv score is 0.663

<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Speechiness']
The cv score is 0.671

<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Speechiness']
The cv score is 0.668

```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Acousticness']
The cv score is 0.659
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Acousticness']
The cv score is 0.677
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Liveness']
The cv score is 0.672
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Liveness']
The cv score is 0.672
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Tempo']
The cv score is 0.66
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Tempo']
The cv score is 0.668
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Duration (ms)']
The cv score is 0.649
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Duration (ms)']
The cv score is 0.649
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Valence']
The cv score is 0.68
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Valence']
The cv score is 0.664
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Speechiness']
The cv score is 0.667
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Speechiness']
The cv score is 0.664
```

```

<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Acousticness']
The cv score is 0.675

<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Acousticness']
The cv score is 0.652

<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Liveness']
The cv score is 0.671

<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Liveness']
The cv score is 0.664

<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Tempo']
The cv score is 0.659

<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Tempo']
The cv score is 0.671

<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Duration (ms)']
The cv score is 0.648

<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Duration (ms)']
The cv score is 0.648

<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Valence']
The cv score is 0.67

<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Valence']
The cv score is 0.658

<class 'numpy.float64'>
Training with columns ['Genre', 'Speechiness', 'Acousticness']
The cv score is 0.673

<class 'numpy.float64'>
Training with columns ['Chord', 'Speechiness', 'Acousticness']
The cv score is 0.676

```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Speechiness', 'Liveness']
The cv score is 0.671

<class 'numpy.float64'>
Training with columns ['Chord', 'Speechiness', 'Liveness']
The cv score is 0.667

<class 'numpy.float64'>
Training with columns ['Genre', 'Speechiness', 'Tempo']
The cv score is 0.662

<class 'numpy.float64'>
Training with columns ['Chord', 'Speechiness', 'Tempo']
The cv score is 0.668

<class 'numpy.float64'>
Training with columns ['Genre', 'Speechiness', 'Duration (ms)']
The cv score is 0.648

<class 'numpy.float64'>
Training with columns ['Chord', 'Speechiness', 'Duration (ms)']
The cv score is 0.648

<class 'numpy.float64'>
Training with columns ['Genre', 'Speechiness', 'Valence']
The cv score is 0.673

<class 'numpy.float64'>
Training with columns ['Chord', 'Speechiness', 'Valence']
The cv score is 0.662

<class 'numpy.float64'>
Training with columns ['Genre', 'Acousticness', 'Liveness']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Chord', 'Acousticness', 'Liveness']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Genre', 'Acousticness', 'Tempo']
The cv score is 0.665

<class 'numpy.float64'>
Training with columns ['Chord', 'Acousticness', 'Tempo']
The cv score is 0.669
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Acousticness', 'Duration (ms)']
The cv score is 0.648
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Acousticness', 'Duration (ms)']
The cv score is 0.648
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Acousticness', 'Valence']
The cv score is 0.662
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Acousticness', 'Valence']
The cv score is 0.679
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Liveness', 'Tempo']
The cv score is 0.66
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Liveness', 'Tempo']
The cv score is 0.669
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Liveness', 'Duration (ms)']
The cv score is 0.648
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Liveness', 'Duration (ms)']
The cv score is 0.648
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Liveness', 'Valence']
The cv score is 0.678
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Liveness', 'Valence']
The cv score is 0.668
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Tempo', 'Duration (ms)']
The cv score is 0.648
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Tempo', 'Duration (ms)']
The cv score is 0.648
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Tempo', 'Valence']
The cv score is 0.662
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Tempo', 'Valence']
The cv score is 0.668
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Duration (ms)', 'Valence']
The cv score is 0.648
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Duration (ms)', 'Valence']
The cv score is 0.648
```

```
<class 'numpy.float64'>
```

```
[65]: display_best_cols(KNN_score)
```

The combination with the highest score is ['Genre', 'Number of Times Charted', 'Popularity'] with a cv score of 0.784.

```
[66]: best_combinations_KNN = ['Genre', 'Number of Times Charted', 'Popularity']

from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train_scaled = X_train.copy()[best_combinations_KNN]

X_test_scaled = X_test.copy()[best_combinations_KNN]
```

```
[67]: from sklearn.neighbors import KNeighborsClassifier

knn = KNeighborsClassifier(n_neighbors = 16)
knn.fit(X_train_scaled, y_train)

#display the model score
knn_train_score = knn.score(X_train_scaled, y_train)
knn_test_score = knn.score(X_test_scaled, y_test)

print(f"KNN Train Score: = {knn_train_score}")
print(f"KNN Test Score: = {knn_test_score}")
```

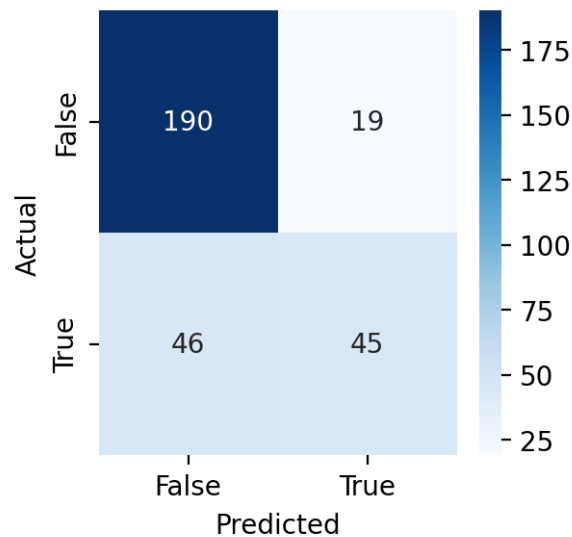
```
KNN Train Score: = 0.8008333333333333
KNN Test Score: = 0.7833333333333333
```




```
[68]: from sklearn.metrics import classification_report
y_pred = knn.predict(X_test_scaled)
report = classification_report(y_test, y_pred)
print(report)


confusion_mat = pd.crosstab(y_test, y_pred, rownames=['Actual'],
    ↳ colnames=['Predicted'])
plt.figure(figsize=(3, 3), dpi=100)
sns.heatmap(confusion_mat, annot=True, cmap='Blues', fmt='g')
plt.show()
```


	precision	recall	f1-score	support
False	0.81	0.91	0.85	209
True	0.70	0.49	0.58	91
accuracy			0.78	300
macro avg	0.75	0.70	0.72	300
weighted avg	0.77	0.78	0.77	300




```
[69]: plot_decision_regions(knn, X_train_scaled, y_train, 'Number of Times Charted',
    ↳ 'Popularity', 'Genre')
```









Analysis

The K-Nearest Neighbors (KNN) model achieved a moderate accuracy of 78% on the testing data. The confusion matrix shows that the model had a precision of 81% and recall of 91% for the False category, indicating that it correctly classified the majority of the non-Top 50 songs. However, the model had a precision of only 70% and recall of 49% for the True category, which represents the Top 50 songs. This indicates that the model still struggled to correctly classify the positive cases.

Looking at the decision regions for the KNN model in the rock and R&B genres, we can observe that the regions are either oddly shaped or do not create any boundaries at all. Specifically look at the sample sizes, although R&B has a lower sample size, it has a higher proportion of True for If Top 50 (13/37). In comparison to rock which only had 16/81 in the Top 50, this is seen in the data from Table 1, Visualization 2. The very few positive data points in these genres, makes it difficult for the model to generalize patterns beyond the specific values it was trained on. As a result, the model may have struggled to define specific decision boundaries in these genres, leading to odd shapes or the absence of boundaries altogether.

The Not Top 50 regions in rap and pop have diagonal lines that differ from the linear regression. However, these decision boundaries make sense in the context of how the categories work together to predict if a song made it into the Top 50. As previously mentioned, we know that a high number of streams and times charted make it more likely that a song will be in the Top 50.

In the case of rap and pop, we can see that there is a diagonal line that divides the Not Top 50 region, which is likely capturing an interaction between streams and number of times charted. Specifically, it appears that when a song has a high number of streams, it is more likely to be in the Top 50, regardless of the number of times it has been charted. However, when a song has a lower number of streams, it becomes increasingly important to have been charted more frequently to be in the Top 50.

6.0.9 C-Support Vector Classification (SVM)

```
[70]: # define a function to compute the score for Support Vector Classification
      ↪ model(sum_model)
      from sklearn import svm

      def score_SVM(cols):
          """
          The function compute the score for sum model.

          Args:
              cols: a list of column names to use as features in the sum model.

          Returns:
              A float representing the accuracy score of the sum model.
          """
          svc = svm.SVC()
          #     svc.fit(X_train[cols], y_train)
          return cross_val_score(svc, X_train[cols], y_train, cv = 10).mean()
```

```
[71]: # observe sum model score
      SVM_score = get_cv_score(score_SVM)
```

```
Training with columns ['Genre', 'Number of Times Charted', 'Streams']
The cv score is 0.719
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Streams']
The cv score is 0.719
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Artist Followers']
The cv score is 0.676
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Artist Followers']
The cv score is 0.675
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Popularity']
The cv score is 0.754
```

```

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Popularity']
The cv score is 0.754

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Danceability']
The cv score is 0.76

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Danceability']
The cv score is 0.759

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Energy']
The cv score is 0.759

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Energy']
The cv score is 0.759

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Loudness']
The cv score is 0.762

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Loudness']
The cv score is 0.762

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Speechiness']
The cv score is 0.759

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Speechiness']
The cv score is 0.759

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Acousticness']
The cv score is 0.759

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Acousticness']
The cv score is 0.759

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Liveness']
The cv score is 0.76

```

```

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Liveness']
The cv score is 0.759

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Tempo']
The cv score is 0.749

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Tempo']
The cv score is 0.75

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Number of Times Charted', 'Valence']
The cv score is 0.759

<class 'numpy.float64'>
Training with columns ['Chord', 'Number of Times Charted', 'Valence']
The cv score is 0.759

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Artist Followers']
The cv score is 0.701

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Artist Followers']
The cv score is 0.701

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Popularity']
The cv score is 0.719

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Popularity']
The cv score is 0.719

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Danceability']
The cv score is 0.719

```



```
<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Danceability']
The cv score is 0.719

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Energy']
The cv score is 0.719

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Energy']
The cv score is 0.719

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Loudness']
The cv score is 0.719

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Loudness']
The cv score is 0.719

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Speechiness']
The cv score is 0.719

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Speechiness']
The cv score is 0.719

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Acousticness']
The cv score is 0.719

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Acousticness']
The cv score is 0.719

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Liveness']
The cv score is 0.719

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Liveness']
The cv score is 0.719

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Tempo']
The cv score is 0.719
```

```

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Tempo']
The cv score is 0.719

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Duration (ms)']
The cv score is 0.72

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Duration (ms)']
The cv score is 0.72

<class 'numpy.float64'>
Training with columns ['Genre', 'Streams', 'Valence']
The cv score is 0.719

<class 'numpy.float64'>
Training with columns ['Chord', 'Streams', 'Valence']
The cv score is 0.719

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Popularity']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Popularity']
The cv score is 0.675

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Danceability']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Danceability']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Energy']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Energy']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Loudness']
The cv score is 0.676

```

```

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Loudness']
The cv score is 0.675

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Speechiness']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Speechiness']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Acousticness']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Acousticness']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Liveness']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Liveness']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Tempo']
The cv score is 0.675

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Tempo']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Duration (ms)']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Duration (ms)']
The cv score is 0.676

<class 'numpy.float64'>
Training with columns ['Genre', 'Artist Followers', 'Valence']
The cv score is 0.676

```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Artist Followers', 'Valence']
The cv score is 0.676
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Danceability']
The cv score is 0.68
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Danceability']
The cv score is 0.68
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Energy']
The cv score is 0.68
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Energy']
The cv score is 0.681
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Loudness']
The cv score is 0.68
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Loudness']
The cv score is 0.68
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Speechiness']
The cv score is 0.68
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Speechiness']
The cv score is 0.68
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Acousticness']
The cv score is 0.68
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Acousticness']
The cv score is 0.68
```

```
<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Liveness']
The cv score is 0.68
```

```

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Liveness']
The cv score is 0.68

<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Popularity', 'Valence']
The cv score is 0.68

<class 'numpy.float64'>
Training with columns ['Chord', 'Popularity', 'Valence']
The cv score is 0.681

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Energy']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Energy']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Loudness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Loudness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Speechiness']
The cv score is 0.678

```

```

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Speechiness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Acousticness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Acousticness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Liveness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Liveness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Danceability', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Danceability', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Loudness']
The cv score is 0.678

```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Loudness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Speechiness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Speechiness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Acousticness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Acousticness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Liveness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Liveness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Energy', 'Valence']
The cv score is 0.678
```

```
<class 'numpy.float64'>
Training with columns ['Chord', 'Energy', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Speechiness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Speechiness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Acousticness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Acousticness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Liveness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Liveness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Loudness', 'Valence']
The cv score is 0.678
```



```

<class 'numpy.float64'>
Training with columns ['Chord', 'Loudness', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Speechiness', 'Acousticness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Speechiness', 'Acousticness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Speechiness', 'Liveness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Speechiness', 'Liveness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Speechiness', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Speechiness', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Speechiness', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Speechiness', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Speechiness', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Speechiness', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Acousticness', 'Liveness']
The cv score is 0.678

```

```

<class 'numpy.float64'>
Training with columns ['Chord', 'Acousticness', 'Liveness']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Acousticness', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Acousticness', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Acousticness', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Acousticness', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Acousticness', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Acousticness', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Liveness', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Liveness', 'Tempo']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Liveness', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Liveness', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Liveness', 'Valence']
The cv score is 0.678

```

```

<class 'numpy.float64'>
Training with columns ['Chord', 'Liveness', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Tempo', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Tempo', 'Duration (ms)']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Tempo', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Tempo', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Genre', 'Duration (ms)', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>
Training with columns ['Chord', 'Duration (ms)', 'Valence']
The cv score is 0.678

<class 'numpy.float64'>

```

```
[72]: display_best_cols(SVM_score)
```

The combination with the highest score is ['Genre', 'Number of Times Charted', 'Loudness'] with a cv score of 0.762.

After conducting a thorough analysis, the combination of features that yielded the highest cross-validation score around of 0.719 includes ['Chord', 'Number of Times Charted', and 'Loudness']. This selection of features aligns with the requirements of utilizing one qualitative feature and two quantitative features for our models.

We can create models that accurately represent the interaction between qualitative and quantitative aspects that influence a song's success by using this feature selection technique. This then makes it possible to make predictions and gain insights into the underlying causes of trends in the music industry that are more precise.

```
[73]: # Define the hyperparameter grid
param_grid = {
    'kernel': ['linear', 'poly', 'rbf', 'sigmoid'],

```

```

    'C': [0.1] + [i for i in range(1, 100, 10)]
}

# Instantiate a svm model with the desired settings
svc_gridS = svm.SVC()

# Instantiate the RandomizedSearchCV object with the svm model and
↳hyperparameter grid
random_search = RandomizedSearchCV(svc_gridS, param_grid, n_iter=20, cv=5,
↳random_state=42)

# Fit the RandomizedSearchCV object to the training data, optimizing for the
↳specified metrics
random_search.fit(X_train[['Chord', 'Number of Times Charted', 'Loudness']],
↳y_train)

# Print the best hyperparameters found by RandomizedSearchCV
print(random_search.best_params_)

```

```
{'kernel': 'rbf', 'C': 71}
```

```

[74]: best_combinations_svm = ['Genre', 'Number of Times Charted', 'Loudness']
svc = svm.SVC(kernel='rbf', C = 71)
svc.fit(X_train[best_combinations_svm], y_train)
best_score_svm = cross_val_score(svc, X_train[best_combinations_svm], y_train,
↳cv = 5).mean()

print(f"The cv score of parameters {random_search.best_params_} and
↳combinations {best_combinations_svm} is {best_score_svm}.")

```

The cv score of parameters {'kernel': 'rbf', 'C': 71} and combinations ['Genre', 'Number of Times Charted', 'Loudness'] is 0.7666666666666667.

6.0.10 Apply the best combination into test set of SVM model

```

[75]: # create Random Forest model with best parameters for our data sets
svc = svm.SVC(kernel='rbf', C = 71)

# fit the model
svc.fit(X_train[best_combinations_svm], y_train)

#display the model score
svm_train_score = svc.score(X_train[best_combinations_svm], y_train)

svm_test_score = svc.score(X_test[best_combinations_svm], y_test)

print(f"SVM CV Score: = {best_score_svm}")

```


```
print(f"SVM Train Score: = {svm_train_score}")
print(f"SVM Test Score: = {svm_test_score}")
```

SVM CV Score: = 0.7666666666666667
SVM Train Score: = 0.7691666666666667
SVM Test Score: = 0.7766666666666666


```
[76]: y_pred = svc.predict (X_test[best_combinations_svm])


report = classification_report(y_test, y_pred)
print(report)
confusion_mat = pd.crosstab(y_test, y_pred, rownames=['Actual'], colnames=
↳ ['Predicted'])
plt.figure (figsize=(3, 3), dpi=100)
sns.heatmap(confusion_mat, annot=True, cmap='Blues', fmt='g')
plt.show()
```


	precision	recall	f1-score	support
False	0.79	0.93	0.85	209
True	0.72	0.43	0.54	91
accuracy			0.78	300
macro avg	0.76	0.68	0.70	300
weighted avg	0.77	0.78	0.76	300




```
[77]: plot_decision_regions(svc, X_test, y_test, 'Number of Times Charted',  
↪ 'Loudness', 'Genre')
```









Analysis

The SVM model achieved an accuracy score of 0.78 on the test data, which suggests that it was able to predict the class labels correctly for 78% of the samples. However, a deeper analysis of the confusion matrix reveals that the model had difficulty predicting the True class label, with a recall score of only 0.43. This indicates that the model has a tendency to misclassify positive samples as negative.

Additionally, we found the selection of Loudness to be a surprising variable to be selected as the highest performing combination, given our earlier EDA analysis that qualities of the song itself (Loudness being one of them) did not typically show a strong correlation with the success of the song.

The decision regions are ultimately less hyper specific when compared to RF, but looking at the individual plots, it appears that the only shape that relatively captured the trends of the data is the graph for Pop. Even for rap, which because is a larger category, has a decision region that does not accurately create general boundaries, exemplified by the two peaks for Not Top 50 in the middle, showing that the model did not generalize the trends in the data as well and rather focused on specific points.

Similar to previous models, the decision regions for rock and r&b show multiple mistakes and specificity that indicates the model does not perform well for these genres which contain minimal

data to create patterns off of.

7 Part 6: Discussion

In this project, our group worked with a set of data containing statistics about songs that made it in the top 200 of Spotify's streaming data to ultimately create a collection of tools and models that could be used to accurately predict if a song will make it into a list of the Top 50 songs streamed. According to a combination of the decision region plots, confusion matrix and cross validation scores, we decided that we would **most recommend using the Random Forest model to predict a song's success** as compared to our other models of linear regression, KNN, and SVM.

Across all 3 measures (i.e. `cv_score`, `train_score` and `test_score`), the Random Forest model consistently yielded the highest score when compared to the other models. The scores are:

- `cv_score`: **0.79** - `'test_score'`: **0.83** - `'train_score'`: **0.93**

It is also the only model which had an accuracy `test_score` of more than 80%. It also produced a precision score on the test data set of 0.84 for False and 0.79 for True, which is the highest precision score among all the other models. From our tuning of the hyperparameters, we concluded that the best combination of variables to run the Random Forest model on is:

- Quantitative Variable 1: **Streams** - Quantitative Variable 2: **Number of Times Charted** - Qualitative Variable: **Genre** This concurred with one of our initial hypotheses that we gathered from just exploring and visualizing the data that **the success of a song making it to the top 50 had less to do with intrinsic qualities** about the song itself, such as Valence, Duration, etc **but rather about the statistics of how the song was consumed/listened to** (such as streams, times charted) or about statistics that would show how the song could reach a larger audience and thus gain more popularity (from an artist with high followers, popularity). It also appears that songs that made it to the Top 50 had a significantly high longevity period. This can be seen throughout the project when we observe that **Number of Times Charted** was a **powerful predictor** of a song's presence in the Top 50 chart.

Even though our Random Forest model produced the best scores, it is important to highlight that it is the only model which yielded a lower `test_score` (0.83) than `train_score` (0.93). While this difference is not too big, it does tell us that overfitting may be slightly present. Moreover, the 0.83 `test_score` is relatively high, which suggests that the model is performing well on unseen data and is a good indication that the model is not overfitting. However, the `test_score` that it yielded was still by far the highest score. We also took into account the fact that our cross-validation score of 0.79 suggests that the model is not overfitting significantly, as it is able to generalize well to new, unseen data. The fact that the cross-validation score is relatively close to the test score is also a strong indication that the model is not extremely overfitting.

In comparison to KNN, LR, and SVM, the RF model emerges as the best choice for dataset. Firstly, Random Forest can capture nonlinear relationships between features and the target variable where models like LR could not achieve. As seen in our first Visualization, there were not strong linear correlations between any of the data points and whether the song made it in the Top 50; therefore, the LR often had to create too large of boundaries and would classify songs in the negative class with far too much frequency and sacrifice its ability to accurately identify when it belonged in the positive.

For the KNN and SVM models, although they were able to create more flexible boundaries in comparison to LR, this does not necessarily mean that they were able to classify data with that

much more accuracy. The decision regions for SVM were still too focused on capturing the behavior of outliers as seen by some of the shapes of the regions in “blobs” that did not make sense in the context of typical patterns we have seen. KNN’s boundaries for R&B and Rock were nearly unusable/did not correctly classify songs within even the test set of data; the graphs for pop and rap showed a slightly clearer boundary, but ultimately were still too complicated and had poor predictive accuracy.

The random forest model emerges as the best choice in terms of decision region performance in comparison to these. One of the key advantages is that it created usable and helpful graphs across all genres, which is something that almost all the other models struggled with, due to the small sample size of R&B and Rock. Another advantage is that the decision regions are able to capture an overall better balance between smoothness and complexity that allows it to effectively capture underlying patterns in the data. Furthermore, the clear boundaries ensure that the model is less sensitive to noise and outliers. This resilience allows it to maintain good performance even in the presence of noisy data or extreme values.

Although there are still many inherent challenges within the RF model, it provides a very solid basis for how we could construct a more detailed machine learning model in the future, especially if we are given access to more data. Past the model itself, the data we were working with also brought on additional difficulties, one of the main ones being the genre category. We had to subjectively sift through the genres manually because due to how the data was given, there was no function or program we could call upon to qualify a song as either “Rap” or “R&B”. We tried to minimize as many errors as possible in this step by only having one individual in charge of the encoding throughout the project, so at least our model was trained and tested on data that fit the same categories. Additionally, there is an inherent unequal classification of genres in this category, as pop songs are the most likely to be the most popular and in the Top 50. Nevertheless, we still found it incredibly interesting to analyze how genres could play a role in learning such statistics about a song, and is an example of how qualitative variables can make the machine learning process more difficult, and how this project could be revisited and improved upon in the future to be able to include more genres.